

Does a duplex-aware CNN add to CLIP prediction after expression control? A bounded benchmark of non-conserved microRNA targets

September 2026

Abstract

Primary question and endpoint: does the CNN improve AGO-CLIP classification beyond seed-site counts, UTR length, and HEK293 expression? Under gene-grouped cross-validation on 435,394 expression-matched (miRNA, gene) rows, pooled AUROC rises from 0.69417 to 0.69494 ($\Delta = +0.00077$). Expression alone reaches 0.68212, exceeding the site-count and CNN models without expression. This conditional estimate is small and does not establish practical utility; a pooled-effect confidence crossed family-and-gene fixed-prediction bootstrap interval includes zero ($[-0.00157, +0.00273]$); a refit-aware crossed interval with fold-training-only scaling also includes zero ($[-0.00010, +0.00295]$); independent-library replication remains outstanding. A second-resource TarBase AGO-CLIP label test passes its locked within-resource positive-delta rule ($+0.01615$, fixed-prediction gene-bootstrap 99% CI $[+0.01459, +0.01773]$), but shared underlying CLIP libraries may link the resources and the median per-miRNA increment is uncertain. The earlier site-count/length-only result (Δ AUROC $+0.0060$ on 476,369 rows) is a secondary, less controlled comparator, not the headline. The model was trained on 264,563 TargetScan conserved sites and scanned non-conserved seed matches, but neither the known miR-7-5p-CDR1as positive control nor the candidate list independently validates new targets. On literature labels, miRDB beats our CNN (miRTarBase pooled AUROC 0.787 vs 0.685), and TarBase reproduces the gap. **What this is not:** a clinical test, a validated target-discovery assay, or evidence that a pairing-map branch is worth its $20\times$ training cost. The resource ledger contains record partitions rather than 317 independent studies; the tool gate remains unverified. The locked random-weight falsifier failed its gate: untrained networks reproduce the apparent dosage depletion. The earlier depletion remains an observed property of the enumerated candidates, not a learned-model discovery. A five-family-held-out transfer test also fails the locked ridge-comparison criterion (mean Pearson $r = 0.705$ for CNN vs 0.726 for ridge+context); there is no independent justification for the pairing-map cost.

Contents

1	Introduction	4
2	Data	4
3	Methods	4
3.1	DuplexCNN architecture	4
3.2	Non-conserved scan	4
3.3	CLIP falsification design	4

4	Results	5
4.1	Benchmark	5
4.2	Discovery scan	5
4.3	CLIP falsification	5
4.4	Pair-map ablation (architecture honesty)	5
4.5	Per-miRNA feature audit: the univariate CNN signal is site-count confounding	6
4.6	Breadth audit: the conditional CNN gain is distributed, not carried by a few miRNAs	6
4.7	Expression confound: HEK293 abundance explains most of the CLIP signal (negative)	7
4.8	Decision rule for current and future use of the CNN score	8
4.9	Cluster-level uncertainty on the expression-adjusted effect	8
4.10	Refit-aware crossed interval with fold-training-only scaling	10
4.11	Pre-registered second-resource AGO-CLIP label test	10
4.12	Biophysical sanity check: duplex MFE vs shuffled controls	11
4.13	Independent literature benchmark: miRTarBase 10.0	12
4.14	Head-to-head vs a public leader: miRDB v6.0 wins (negative)	12
4.15	Second independent benchmark: TarBase v9.0 (negative for the CNN)	12
4.16	Third arm: pure duplex thermodynamics (RNAhybrid) is near chance	13
4.17	Fourth arm: alignment-based biophysics (miRanda) sits between	13
4.18	Functional coherence audit: g:Profiler finds no pathway signal in predicted or CLIP target sets (negative)	14
4.19	Network coherence audit: STRING edges in CLIP target sets are explained by expression (negative)	15
4.20	miRNA abundance does not predict which miRNAs the CNN ranks well (ENCODE; negative)	16
4.21	Tissue target avoidance (GTEx): conserved targets avoid the miRNA's tissue, CNN non-conserved targets do not	16
4.22	Reverse lookup (Enrichr): CNN top-200 sets name their own miRNA beyond expression, but the site-count win does not replicate	18
4.23	Reactome pathway audit with controls fixed in advance: two gates failed, verdicts inconclusive	19
4.24	Input integrity: every miRNA sequence and seed used here matches RNACentral	19
4.25	Independent curated labels (OmniPath, miRTarBase removed): CNN above chance, but not above site count or expression	20
4.26	Transcription-regulator enrichment (UniProt): present in conserved targets, absent in CNN non-conserved targets	20
4.27	Loss-of-function constraint (gnomAD): conserved targets are constrained, CNN candidates are the least constrained set	21
4.28	Independent dosage label (ClinGen): the depletion beyond UTR length does not replicate	22
4.29	Third dosage label (Collins et al. 2022 rCNV scores): the depletion survives UTR-length matching	23
4.30	Mendelian disease genes (HPO): depletion explained by UTR length	24

4.31	ClinVar pathogenic-variant genes: the positive-control gate failed	24
4.32	Cell-essential genes (DepMap): no depletion, and the label is uninformative for conserved targets	25
4.33	Protein-complex subunits (Complex Portal): the dosage-balance prediction fails	25
4.34	Study bias (NCBI Gene, PubMed counts): it explains the curated-label depletion but not the LOEUF depletion	26
4.35	DECIPHER haploinsufficiency predictions: depletion beyond expression and study bias, not beyond UTR length	26
4.36	GeneBayes s_{het} : population-based labels agree with each other, curated and model-based labels do not	27
4.37	Paralog buffering (Ensembl BioMart): no paralog excess; both gates failed	27
4.38	Gene families (HGNC): the depletion is not a zinc-finger artefact; coverage gate failed . . .	28
4.39	Orphanet rare-disease genes: depletion vs expression-matched only; gone against publication-matched sets	28
4.40	Mouse mutant phenotypes (MGI): weak, mixed support	28
5	Independent benchmark audit: every evaluated miRNA	29
5.1	Interpretation of the independent ranking	29
6	Documented negatives	31
6.1	Locked random-weight and score-permutation falsifier	32
6.2	Family-held-out transfer test and the pair-map decision	33
6.3	Fixed-width absolute UTR-length-bin sensitivity audit	33
7	Falsification paths	34
8	Mathematical derivations	36
8.1	Convolutional duplex scoring	36
8.2	Pearson and Spearman benchmarks	36
8.3	Fisher exact test for CLIP enrichment	36
8.4	Odds ratio	37
8.5	AUROC as a Mann–Whitney statistic	37
8.6	Bootstrap confidence interval for ΔAUROC	37
8.7	Permutation null	37
8.8	Seed-match enumerator	37
8.9	Wilcoxon signed-rank for paired per-gene comparisons	38

1 Introduction

Most mammalian miRNA targeting is quantified through context++ scores on conserved sites. Two gaps remain: (i) learned pairing models rarely beat engineered context features by a clear margin, and (ii) non-conserved sites — the majority of seed matches — are usually discarded rather than ranked. The single primary question is whether the CNN adds CLIP discrimination conditional on expression, seed-site counts, and UTR length. Other tests are secondary stress tests and cannot rescue a failed primary endpoint.

2 Data

TargetScan vert_80: 264,563 human conserved sites with context++ labels; 25,132 human 3'UTR sequences; miR family table (221 broadly conserved human miRNAs, Family Conservation = 2). ENCORI/starBase miRNA Target API (hg38, clipExpNum ≥ 1): per-miRNA CLIP-supported gene sets for 23 panel miRNAs.

3 Methods

3.1 DuplexCNN architecture

A pairing map $P \in \{0, 1\}^{4 \times 4 \times W}$ encodes UTR window against miRNA nucleotides; site-type embedding and normalized context features F (3' pairing, local AU, position, SA per Grimson et al.) join at the head:

$$\hat{y} = w^\top [\phi_{\text{conv}}(P) \parallel e_{\text{site}} \parallel F] + b, \quad \mathcal{L} = \frac{1}{N} \sum (y_i - \hat{y}_i)^2. \quad (1)$$

Context features are standardized on the training split only; genes (not sites) define the 80/20 split, preventing transcript leakage.

3.2 Non-conserved scan

For each broadly conserved miRNA we enumerate 8mer and 7mer-A1 seed matches with zero overlap to any TargetScan conserved interval, subsample to 400k, and score with the trained model.

3.3 CLIP falsification design

Unit of analysis: (miRNA, gene). Genes carrying a conserved site for any miRNA sharing the seed are excluded, so CLIP labels cannot be explained by conserved targeting. Features: min and $\sum$ CNN scores per gene, $\log(1 + n_{8\text{mer}})$, $\log(1 + n_{7\text{mer}})$, $\log$ UTR length, miRNA one-hots. Logistic models (base vs base+CNN) compared by 5-fold CV AUROC, bootstrap CI on Δ AUROC ($B = 1000$), and a within-miRNA permutation null (20 shuffles).




Figure 1: Held-out-gene Pearson r by model. Sequence-only CNN v1 (0.669) merely matched ridge (0.664) — archived as a non-win; adding context features gives DuplexCNN 0.811 vs ridge+context 0.779.

4 Results

4.1 Benchmark

4.2 Discovery scan

1,453,439 novel non-conserved seed matches; 400k scored. Top overall: hsa-miR-7-5p→CDR1as ($\hat{y} = -1.57$), the ciRS-7 sponge — rediscovered with no conservation information. Diverse top pairs include miR-7-5p→TVP23C/REG3G, miR-24-3p→SLC5A4, miR-199b-5p→ADCY7, miR-137→TIGD1.

4.3 CLIP falsification

Pooled (23-miRNA panel): AUROC 0.5905 (covariates) → 0.5979 (+CNN); $\Delta = +0.0075$, CI [0.0057, 0.0093], permutation $p = 0.048$. **Extended-panel replication (132 miRNAs, 476,369 (miRNA, gene) pairs):** AUROC 0.6047 (covariates) → 0.6108 (+CNN); $\Delta = +0.0060$, bootstrap CI [+0.0053, +0.0067], within-miRNA permutation $p = 0.048$ (20 shuffles); per-miRNA AUROC of the negated CNN-min score median 0.531, 85% of miRNAs above chance (sign convention audited below). This secondary no-expression comparison reproduces a small effect on a larger panel, but its positive verdict does not carry over to the expression-controlled primary endpoint. **Verdict: this is not the primary endpoint; expression confounding is assessed below.** The discovery list is CLIP-supported above chance but the CNN is not a stand-alone oracle for non-conserved targeting; candidates are released as ranked hypotheses, not validated targets.

4.4 Pair-map ablation (architecture honesty)

The explicit complementarity pair-map branch is ablated under the identical split, seed, and epochs: full DuplexCNN Pearson 0.8121 / Spearman 0.7928 vs no-pair-map 0.8092 / 0.7898 ($\Delta r = +0.0029$, RMSE




Figure 2: Score curve over the top 100 diverse (miRNA, gene) pairs.

−0.0005) — but at 333.5s vs 16.3s train time (20× cheaper without it). Verdict: the pair-map is a minor contributor for context++ prediction; the sequence towers carry nearly all signal. The branch is retained in the archived trained artifact to permit reproduction, but is deprecated as a default for new model training. The family-held-out test below did not earn an exception at its 20× cost.

4.5 Per-miRNA feature audit: the univariate CNN signal is site-count confounding

To ask where the pooled Δ AUROC lives, we computed univariate AUROCs per miRNA for every feature in the CLIP model (132 miRNAs, 476,369 pairs, `results/per_mirna_feature_auroc.json`). This audit surfaced a sign-convention finding: the per-miRNA AUROC in the falsification runs is computed on the *negated* min-score (`roc_auc_score(y, -min)`), so the raw CNN min-score is univariately *anticorrelated* with CLIP labels (raw median AUROC 0.469 vs negated 0.531; the two are exact complements, verified per miRNA). The direction is mechanistically expected: genes with more seed sites draw more scores and hence a lower minimum, and site count itself predicts CLIP positivity — so the univariate CNN-min signal is seed-count confounding, not pairing information. Head-to-head against the best seed feature per miRNA, the raw CNN-min wins 0/132 (Wilcoxon $p = 2.1 \times 10^{-23}$). The honest reading of the pooled result is therefore narrow: conditioned on site counts and UTR length the CNN adds a small complementary signal (Δ AUROC +0.0060), but it carries no usable standalone ranking signal within a miRNA. Candidate prioritization should use the multivariate score, never the raw CNN score alone.

4.6 Breadth audit: the conditional CNN gain is distributed, not carried by a few miRNAs

The pooled Δ AUROC of +0.0060 could in principle be driven by a handful of miRNAs. We decomposed it: for each of the 131 panel miRNAs with ≥ 50 CLIP-positive and ≥ 50 negative genes we re-fit both logistic models (covariates $[n_8, n_7, \text{len}]$ vs covariates+CNN $[\text{min}, \text{sum}]$) under 5-fold cross-validation within that miRNA alone and recorded the per-miRNA Δ AUROC (`scripts/clip_per_mirna_delta.py`, `results/clip_per_mirna_delta.json`, three hermetic tests). The gain is broad but small: median +0.0046, IQR $[-0.0013, +0.0155]$, 64.9% of miRNAs positive, Wilcoxon signed-rank $p = 2.0 \times 10^{-9}$ across miRNAs (Figure 5). The largest gains are real but modest (hsa-miR-500a-3p +0.084, hsa-miR-302a-




Figure 3: Per-miRNA AUROC of the negated CNN min-score against ENCORI CLIP support. Dashed: chance. Sign convention audited in the per-miRNA feature audit.

3p +0.079, hsa-miR-93-5p +0.072, hsa-miR-20a-5p +0.070; three of the five share the miR-17-family seed and all five have large pair counts), and the worst cases are only mildly negative (hsa-miR-615-3p −0.014, hsa-miR-652-3p −0.014). Two honest readings follow. First, the pooled effect is not an artifact of one dominant miRNA: the sign test across 131 independent miRNAs would be significant on its own. Second, the per-miRNA conditional gain is an order of magnitude smaller than the headline pooled number for most of the panel, so for any individual miRNA the CNN score remains a weak assistant to site-count features, not a replacement — consistent with the univariate audit above.

4.7 Expression confound: HEK293 abundance explains most of the CLIP signal (negative)

ENCORI AGO-CLIP libraries come mostly from HEK293 cells, and a transcript that is abundant in HEK293 yields more CLIP reads whether or not the miRNA regulates it. We tested this directly with the Human Protein Atlas cell-line RNA atlas (`rna_celline.tsv.zip`, HEK293 nTPM; derived table committed as `results/hpa_hek293_ntpm.tsv`). Of the 476,369 committed CLIP rows, 435,394 matched an HPA gene symbol (1,891 genes unmatched and dropped). CLIP-positive genes are far more expressed: median $\log(1 + \text{nTPM})$ 2.79 vs 1.76 for negatives (Mann–Whitney p below double-precision underflow). Under gene-grouped 5-fold cross-validation (no gene in both train and test; `scripts/expression_confound.py`, `results/expression_confound.json`, four hermetic tests) the pooled AUROCs (Figure 6) are: site counts $[n_8, n_7, \text{len}]$ 0.568; site counts+CNN 0.573; HEK293 expression alone 0.682; site counts+expression 0.694; and site counts+expression+CNN 0.695. Expression alone beats every sequence model we have. Adding expression shrinks the CNN’s conditional pooled gain from +0.0048 to +0.0008, about a six-fold reduction. Per miRNA (130 miRNAs with ≥ 50 positives and negatives, covariates now including expression) the gain survives in sign but is tiny: median +0.0004, IQR $[-0.0007, +0.0029]$, 60.8% positive, Wilcoxon $p = 5.1 \times 10^{-4}$. *Primary-endpoint verdict*: the point estimate is positive but only +0.0008 AUROC. The later crossed family-and-gene fixed-prediction interval includes zero, so no confident positive pooled effect, practical-value or target-discovery pass is declared. The earlier +0.0060 comparison is sec-




Figure 4: Per-miRNA univariate AUROC: raw CNN min-score vs the best seed-count/length feature. Every one of 132 miRNAs sits below the diagonal.

ondary after controlling the dominant confound. Every earlier CLIP number in this paper should be read with this correction. The obvious next control, CLIP read depth normalized by RNA-seq in the same library, needs per-library data that ENCORI does not expose.

4.8 Decision rule for current and future use of the CNN score

The expression-controlled primary comparison gives a pooled $+0.00077$ AUROC point estimate. A crossed family-and-gene fixed-prediction interval below includes zero, so even the sign is uncertain under that sensitivity test; a refit-aware interval under fold-training-only scaling, reported below, also includes zero. Independent-library replication remains open. A second-resource TarBase test passes its locked within-resource rule, but its source libraries can overlap ENCORI and its per-miRNA median interval includes zero; it does not establish the independent-library condition. The full CNN also failed the locked family-held-out ridge comparison reported below. Our current rule is therefore *do not assert practical added value*: use site counts, UTR length and measured HEK293 expression as the simpler comparator for this dataset, and do not rank targets by raw CNN min-score alone. A future claim that the CNN adds value requires a separately locked, externally replicated comparison on matched library-normalized labels, a dependence-aware interval on its conditional pooled AUROC increment, and a prospectively defined minimum useful increment for the intended decision. No such practical threshold has been supplied here; selecting one after seeing a $+0.00077$ point estimate would make a misleading pass. This is an operational interpretation of limited evidence, not a new statistical gate or a recommendation to use the comparator clinically.

4.9 Cluster-level uncertainty on the expression-adjusted effect

A 10,000-replicate percentile bootstrap resampled the 130 miRNAs as clusters, using the stored per-miRNA expression-adjusted conditional AUROC increments (the script and result are committed in the reposi-




Figure 5: Distribution of per-miRNA conditional ΔAUROC across 131 miRNAs. Dashed: zero; dotted: median (+0.0046). 64.9% of miRNAs are positive; Wilcoxon $p = 2.0 \times 10^{-9}$.




Figure 6: Left: pooled AUROC of nested logistic models on CLIP labels under gene-grouped CV. HEK293 expression (HPA nTPM) alone outperforms site counts and the CNN. Right: per-miRNA ΔAUROC of adding the CNN to site counts+expression (130 miRNAs; dotted line = median +0.0004).

tory; random seed 20260927). The across-miRNA median is +0.00042, with a bootstrap 95% interval [+0.00009, +0.00099]; 79 of 130 miRNAs have positive increments. This interval describes the median within-miRNA increment, *not* the pooled +0.00077 AUROC and not a confidence interval on the change for a new individual miRNA. Related miRNAs can share seeds and genes; miRNA-level resampling may understate dependence by family or common genes. In an exploratory family-cluster bootstrap, the 130 miRNAs map to 103 TargetScan human families; resampling entire families 10,000 times gives a 95% percentile interval [+0.000023, +0.001218] for the same median increment (`results/family_cluster_interval.json`). This accounts for within-family ties but not shared genes across families, and is still not an interval for the pooled +0.00077 AUROC. Separately, an exploratory fixed-prediction gene-cluster bootstrap resampled 16,835 gene symbols, carrying all 435,394 matched miRNA–gene rows with them; 300 replicates produced a pooled-increment percentile interval of [+0.000332, +0.001183] (0 nonpositive replicates). The original out-of-fold AUROCs were reproduced exactly; details are in `results/pooled_gene_interval.json`. This conditional calculation holds predictions fixed; it neither refits models nor accounts jointly for family and gene dependence. The original global scaling before folds can transmit test-distribution information into

training. Thus the interval is a sensitivity check, not a full pooled uncertainty bound or a repair of preprocessing. A stronger crossed-cluster weighted bootstrap drew entire miRNA families and genes with replacement, assigning each row the product of its family and gene counts. It included 104 families across all matched rows (the earlier median analysis had 103 because it excluded miRNAs with few positive or negative labels), 16,835 genes and 300 seeded draws. The pooled-increment interval was $[-0.001568, +0.002733]$; 70 draws were nonpositive (`results/pooled.twoway_interval.json`). The family-by-gene product weighting is an exploratory crossed-dependence sensitivity analysis, not a formal model-refit interval. Crucially, its interval includes zero: the sign of the tiny primary pooled increment is uncertain once related miRNAs and genes are stressed together. The original global scaling still precedes fold assignment; no interval repairs that preprocessing weakness. Nor does any interval supply independent-library replication or practical value.

4.10 Refit-aware crossed interval with fold-training-only scaling

The crossed interval above still refits nothing: it holds the original out-of-fold predictions fixed and retains the source experiment’s global scaling, which precedes fold assignment. To test whether that preprocessing transduction, rather than dependence alone, drives the uncertainty, a pre-registered evaluator-stack refit (CNN feature extractor frozen; declared cost bound) re-drew the same crossed family-and-gene weighted bootstrap with one change: inside each of the 300 seeded draws, the weighted standardization scalers were fit on the training folds of that draw only and then applied to the held-out fold, so fold-held-out genes never entered scaler statistics. An integrity gate ran the refit pipeline at uniform weights first and reproduced the original global-scaled zero-draw baselines (0.69416990 / 0.69493803) before any replacement draw was observed; an earlier attempt whose scalers had touched fold-held-out genes was stopped at its checkpoint, never read as an outcome, and is retained as an audit trail. The corrected zero-draw baselines are 0.69417050 (expression-only) and 0.69493857 (with conserved-site context), increment $+0.00077$. The refit-aware 95% percentile interval over 300 draws (0 skipped) is $[-0.00010, +0.00295]$, with 14 nonpositive draws: the interval still includes zero, so the sign of the tiny pooled increment remains uncertain under the strongest dependence-and-refit sensitivity test we have run. This does not restore practical utility, and it is not an independent-library replication; features remain TargetScan-trained with labels disjoint from CLIP. What it establishes is that fold-training-only scaling does not rescue the sign: the uncertainty is dependence, not preprocessing transduction. (`results/refit_interval_folds.json`; protocol docs/PREREG_REFIT_INTERVAL_FOLD_SCALING_20260928.md.)

4.11 Pre-registered second-resource AGO-CLIP label test

The dated pre-registration (docs/PREREG_SECOND_CLIP_20260927.md) set a positive pooled AEC-minus-AE increment and a gene-cluster bootstrap interval excluding zero as its replication-support rule, before second-resource outcomes. It also required at least 50 miRNAs with at least 50 positives and 50 negatives after gene filtering. An ENCODE K562 AGO2 eCLIP replicated-peak file (ENCSR048YXG, ENCFF111TWF; 12,307 peaks) yielded only 41 eligible miRNAs after site-to-gene mapping. It is attempted but ineligible, and its outcomes do not enter the inference.

We next used DIANA TarBase v9.0 [Skoufos et al.(2024)], which re-analyses raw CLIP experiments rather than re-exporting ENCORI. Before eligibility or model outcomes, we restricted pair support to HITS-CLIP, PAR-CLIP, CLASH and qCLASH, and copied the original 476,369 non-conserved (miRNA, gene) feature rows verbatim, replacing only their support labels. There are 120 eligible miRNAs on the pre-specified 50/50 criterion. Matching HPA HEK293 expression retains 435,394 rows (118,244 positives, 317,150 negatives; 1,891 gene symbols unmatched). Under the same gene-grouped five-fold logistic comparison, the pooled AUROC is 0.656177 for site counts, UTR length and expression (AE), and 0.672323 with CNN min/sum

added (AEC), a $+0.016146$ increment. The pre-specified gene-cluster percentile bootstrap (1,000 draws, seed 0), resampling all rows for each gene while holding out-of-fold predictions fixed, gives a 95% interval $[+0.014849, +0.017389]$ and a 99% interval $[+0.014586, +0.017726]$; all 1,000 draws are positive. This passes the locked positive-delta and zero-exclusion rule, even at the specified one-sided $\alpha = 0.01$. No cutoff was moved after seeing the result.

A pre-specified secondary analysis fits each eligible miRNA separately: median increment $+0.000094$, interquartile range $[-0.001111, +0.002722]$, 61/120 positive. The 10,000-draw miRNA-cluster bootstrap 95% interval for that median is $[-0.000330, +0.000520]$. Thus pooled discrimination improves on this resource’s labels, but a typical individual-miRNA improvement is not established. TarBase and ENCORI can share underlying CLIP libraries; different processing and curation are *not* independent-library validation. The bootstrap is conditional on fixed predictions, does not refit models or repair global scaling before fold assignment, and omits crossed miRNA-family/gene dependence. It cannot resolve the ENCORI crossed-dependence interval spanning zero, prove practical value, restore failed biological discovery gates, or validate candidate pairs. The committed label audit, analysis script, result JSON and source ledger record eligibility counts, exact outputs, input URLs and checksums. The ENCODE path and why it failed eligibility are recorded in the same ledger.

4.12 Biophysical sanity check: duplex MFE vs shuffled controls

As a sequence-level pipeline validation, the 20 top diverse candidate windows (± 25 nt) were cofolded with their mature miRNA (ViennaRNA RNAcofold) against 10 mononucleotide-shuffled miRNA controls each (results/rnafold_validation.json). Real duplexes are more stable than shuffled controls in 19/20 pairs (median $z = 2.33$; Wilcoxon $p < 10^{-4}$); the exception, miR-7-5p→REG3G ($z = -0.57$), is kept in the candidate table as a flagged row. Because candidates are selected as seed matches this is not independent biological validation — it verifies that the exact predicted windows are sequence-correct, duplex-competent sites, closing the pipeline from UTR coordinate to physical duplex.




Figure 7: RNAcofold MFE for the 20 top candidate windows with the real miRNA vs shuffled-miRNA controls.

4.13 Independent literature benchmark: miRTarBase 10.0

All evaluations above use TargetScan or ENCORI labels. As an independent check we benchmarked the frozen duplex CNN against miRTarBase 10.0 [Huang et al.(2022)]: positives were strong-evidence (reporter-assay) human MTIs for the 132-miRNA panel (miRTarBase_SE.R.csv, 12,981 records), and negatives were genes with no human MTI of any evidence class in the full 337 MB hsa.MTI.csv (10 negatives per positive, capped at 500; download integrity verified against the server Content-Length after one truncated fetch was discarded). Genes are scored by the minimum CNN score over enumerated 8mer/7mer-A1 seed windows (results/mirtarbase_validation.json). On 87 miRNAs with ≥ 10 validated positives (45 skipped), 8,075 rows and 2,164 positives, the pooled AUROC is 0.6829; the median per-miRNA AUROC is 0.6926 with 85/87 miRNAs above chance (within-miRNA z Mann–Whitney $p = 1.0 \times 10^{-152}$; 1000-rep permutation $p \leq 9.99 \times 10^{-4}$, the resolution floor). Because the CLIP audit identified seed-site count as a confound, we re-scored the identical rows with pure site count as the predictor (results/mirtarbase_sitecount_control.json): count alone reaches 0.6573 pooled / 0.6675 median, so most of the literature signal is seed abundance. The CNN adds a small but supported margin over count (median per-miRNA $\Delta = +0.0349$, positive for 65.5% of miRNAs, Wilcoxon $p = 1.2 \times 10^{-3}$). Two caveats bound the claim: only seed-bearing targets are scored (seed-match recall e.g. 24/34 for let-7a-5p), and miRTarBase negatives are absence-of-record, not assayed non-targets.

4.14 Head-to-head vs a public leader: miRDB v6.0 wins (negative)

We compared the frozen duplex CNN against miRDB v6.0 (MirTarget) [Chen and Wang(2020)] on the identical miRTarBase rows (scripts/mirdb_benchmark.py, results/mirdb_benchmark.json). The full miRDB prediction file (59 MB, 3,375,741 human rows) was downloaded; RefSeq accessions were mapped to gene symbols with the MyGene.info batch API [Xin et al.(2016)] (29,460 of 29,509 mapped). A gene’s miRDB score is the maximum over its transcripts; pairs absent from miRDB (which reports only scores ≥ 50) score 0. On the 84 panel miRNAs present in miRDB (7,909 rows, 2,113 positives), miRDB reaches pooled AUROC 0.787 vs 0.685 for our CNN; median per-miRNA AUROC is 0.791 vs 0.693 (median $\Delta = +0.093$ in miRDB’s favour), miRDB wins on 74/84 miRNAs, Wilcoxon $p = 2.4 \times 10^{-12}$. miRDB lists 68.9% of positives but only 15.5% of negatives. *Verdict*: on literature-validated targets our model is clearly behind the public leader. One caveat favours miRDB: MirTarget was trained on external CLIP/expression data that may overlap literature-validated targets, so this is not a leakage-free comparison; but the size and consistency of the gap make it a real negative, and we make no state-of-the-art claim for target ranking.

4.15 Second independent benchmark: TarBase v9.0 (negative for the CNN)

To check that the miRTarBase result is not an artefact of one curation pipeline, we relabelled the *same* 7,909 rows (84 miRNAs, identical gene universe and frozen CNN scores) using TarBase v9.0 [Skoufos et al.(2024)], the DIANA collection of experimentally supported interactions (4,724,537 human records scanned; scripts/tarbase_benchmark.py, results/tarbase_benchmark.json). A row is positive if the pair has at least one TarBase record in the evidence class. Three predictors are compared: our CNN, miRDB v6.0, and the plain seed-site count (8mer + 7mer-A1).

Label class	Pos.	Pooled AUROC			Median per-miRNA AUROC		
		CNN	miRDB	count	CNN	miRDB	count
Any method	3,978	0.609	0.666	0.601	0.583	0.672	0.609
Low-yield (reporter/WB/qPCR)	1,019	0.661	0.754	0.643	0.654	0.766	0.639
AGO-CLIP/chimeric	3,460	0.591	0.644	0.581	0.571	0.668	0.590

miRDB beats the CNN on 73/83 miRNAs for any-method labels (Wilcoxon $p = 3.2 \times 10^{-13}$), 56/69 for low-yield labels ($p = 3.0 \times 10^{-8}$) and 73/80 for CLIP labels ($p = 2.1 \times 10^{-12}$). Against the seed-site count the CNN wins the pooled AUROC by less than 0.02 in every class, but *loses* per miRNA: it beats the count on only 32/83 miRNAs for any-method labels ($p = 0.0078$) and 32/80 for CLIP labels ($p = 0.013$); for low-yield labels it wins 42/69 (median 0.654 vs 0.639, $p = 0.056$, not significant). *Verdict:* the second database reproduces the miRDB gap and removes the small margin over site counting that the miRTarBase control showed. Our CNN adds no reliable per-miRNA value over counting seed matches on TarBase labels.

A data-quality finding: 2,199 of the rows used as miRTarBase *negatives* have TarBase support (2,063 of them from CLIP-class methods, only 7 from low-yield methods). The miRTarBase negative set is therefore contaminated by CLIP-supported interactions, which lowers every predictor’s measured AUROC on that benchmark. Low-yield TarBase positives overlap miRTarBase strong-evidence positives almost fully (1,012 of 1,019), so that label class is not independent of miRTarBase; the any-method and CLIP classes are the more independent tests. Caveats: TarBase CLIP records overlap the ENCORI CLIP data used in our falsification analyses (not in CNN training), and miRDB’s MirTarget was trained on CLIP and array data (leakage in its favour).

4.16 Third arm: pure duplex thermodynamics (RNAhybrid) is near chance

The miRDB comparison has an asymmetry: miRDB’s MirTarget is a learned model with many features, our CNN is sequence-only. We added the cleanest possible biophysical arm: RNAhybrid 2.1.2 minimum-free-energy duplex prediction (`-s 3utr_human`, best hit per transcript, minimum over a gene’s transcripts; no training data, no conservation, no machine learning) on the same committed miRTarBase rows (`script rnahybrid_benchmark.py; results/rnahybrid_benchmark.json; results/rnahybrid_rows.csv`; three hermetic tests). Rows with no RNAhybrid hit are scored one step below the weakest observed hit. On 80 panel miRNAs (7,914 rows) RNAhybrid is near chance: median per-miRNA AUROC 0.542 (pooled 0.522) with median gene coverage 54%. The CNN wins 76/80 miRNAs (medians 0.693 vs 0.542, Wilcoxon $p = 2.0 \times 10^{-14}$). Together with the miRDB arm this gives a clean three-way ordering on identical rows (Figure 8): a full machine-learning predictor (miRDB 0.791) > our sequence-only CNN (0.693) > pure duplex thermodynamics (0.542). Two honest readings. First, the CNN’s margin over the strongest physics-only baseline is large, so the sequence encoding does capture real, generalizing signal beyond hybridization energy. Second, the equally large gap from CNN to miRDB shows that the extra signal miRDB exploits (cross-species conservation, site context, CLIP-trained features) is worth more than our whole model: thermodynamics is not the bottleneck, feature breadth is.

4.17 Fourth arm: alignment-based biophysics (miRanda) sits between

The RNAhybrid arm isolates duplex free energy; miRanda 3.3a is the other canonical pure-biophysics predictor class - Smith-Waterman seed-weighted alignment plus Vienna energy, default cutoffs, no training data and no conservation (`scripts/miranda_benchmark.py; results/miranda_benchmark.json; results/miranda_rows.csv`; hermetic tests). On the identical 80-miRNA, 7,914-row miRTarBase benchmark, miRanda lands strictly between the two existing arms: pooled AUROC 0.610 (median per-miRNA 0.624, median gene coverage 80%) versus CNN 0.682 (median 0.693) and RNAhybrid 0.522. The CNN wins 66/80 miRNAs against miRanda (Wilcoxon $p = 6.0e - 08$). The four-way ordering on this benchmark is now miRDB (0.787) > DuplexCNN (0.682) > miRanda (0.610) > RNAhybrid (0.522) (Figure 9). The reading is consistent with the third arm: adding alignment shape on top of duplex energy buys real signal (+0.088 pooled over RNAhybrid), but the learned sequence encoding buys more (+0.073 over miRanda), and miRDB’s feature breadth buys most.




Figure 8: Per-miRNA AUROC on the committed miRTarBase rows (77 miRNAs present in both benchmarks). Left: distributions for RNAhybrid, DuplexCNN, miRDB. Right: RNAhybrid vs DuplexCNN per miRNA; the CNN wins 76/80.

4.18 Functional coherence audit: g:Profiler finds no pathway signal in predicted or CLIP target sets (negative)

If the CNN ranks real targets, its top-ranked genes for a miRNA might share biology. We tested this with the g:Profiler g:GOST service (version e114_eg62_p19_27110d83; scripts/gprofiler_enrichment.py; results/gprofiler_enrichment.json; hermetic tests) on 13 well-studied miRNAs. For each miRNA the background is the set of genes scored for that miRNA, and we query four gene sets of size $K = 200$: CNN top- K , CLIP-positive genes ranked by CNN score, the top- K genes by HEK293 abundance (confound control) and 3 seeded random draws. Terms are called significant under g:SCS, which corrects the one-sided hypergeometric tail

$$p = \sum_{i \geq k} \frac{\binom{K_i}{i} \binom{N-K_i}{n-i}}{\binom{N}{n}} \quad (2)$$

(query size n , background N , term size K_i , overlap k) for the dependence among GO terms. Gates, fixed before any group comparison, all passed: a 20-gene canonical cell-cycle positive control recovers GO:0007049 (top hit KEGG cell cycle, $p = 3.5 \times 10^{-31}$); the median random set has 0 significant terms; the median unmapped fraction is 6.75%. All three hypotheses are falsified. CNN top- K sets have zero significant terms for 13/13 miRNAs and CLIP-positive sets for 12/13, so neither beats random (H1: 0 wins, 0 losses, $p = 1.000$; H2 term overlap with the CLIP set: 0 wins, 0 losses). The expression control returns 47–115 terms per miRNA and beats the CNN on 13/13 (H3, Figure 10). A post-hoc arm that we did not pre-register swaps to the default whole-genome background: CNN top- K still loses to the first random set (0 wins, 6 losses; median 0 vs 1 terms), consistent with (but not a test of) the CLIP-scored universe being biased toward expressed, well-annotated genes. The reading matches the expression-confound section: at $K = 200$, target lists from sequence or CLIP carry no detectable pathway coherence, while abundance alone carries a lot. Pathway enrichment of predicted miRNA targets should not be read as evidence of targeting unless it beats an expression-matched control.




Figure 9: Left: pooled AUROC on the committed miRTarBase rows for the four predictor arms (miRDB from results/mirdb_benchmark.json; the other three on identical rows, $n = 7,914$). Right: miRanda versus DuplexCNN per miRNA; orange points are miRNAs where miRanda wins.




Figure 10: Significant g:GOST terms per miRNA for four gene sets of 200 against each miRNA's scored-gene background. Expression top-200 (orange) is the confound control.

4.19 Network coherence audit: STRING edges in CLIP target sets are explained by expression (negative)

The g:Profiler section tests term enrichment; STRING v12.0 tests whether a target set is more densely connected than chance (scripts/string_coherence.py; results/string_coherence.json; hermetic tests). Same panel, sets and seeds as above; the statistic is the edge density $\rho = E / \binom{n}{2}$ among mapped query proteins with STRING combined score ≥ 0.7 . Gates passed: the cell-cycle control gives 133 edges (bar 50); median random density 0.00145 (bar 0.01); median mapped fraction 97.5%. CNN top-200 sets are no denser than random (H1: 6 wins, 7 losses, $p = 0.709$, falsified) and lose to the expression top-200 control on 13/13 (H3, falsified). CLIP-positive sets are denser than random on 13/13 ($p = 1.2 \times 10^{-4}$), so pre-registered H2 is confirmed. Because CLIP detection depends on expression, we then added a post-hoc control that we did not pre-register: random sets drawn from each miRNA's universe to match the CLIP set's HEK293 expression-decile histogram. Against it the CLIP excess does not survive (8 wins, 4 losses, one tie excluded, $p = 0.194$; Figure 11). So the one positive in this audit is explained by expression, the same con-

found that dominates the CLIP labels and the g:Profiler terms. We report H2 as confirmed-then-explained, not as evidence that miRNA targets form network modules.




Figure 11: Left: STRING edge density per miRNA for CNN top-200, CLIP+ top-200, seeded random and expression-matched random sets. Right: distribution across miRNAs with the expression top-200 control, which is several-fold denser.

4.20 miRNA abundance does not predict which miRNAs the CNN ranks well (ENCODE; negative)

CLIP labels need AGO loaded with the miRNA, so abundant miRNAs could give cleaner labels and higher per-miRNA AUROC. We measured abundance from 21 ENCODE GRCh38 microRNA-seq quantification files covering 12 cell lines (accessions in `results/encode_mirna_abundance.json`; Ensembl REST lookup/id for gene symbols; `scripts/encode_mirna_abundance.py`; hermetic tests). Per file we take $a_g = \log_2(\text{CPM}_g + 1)$, sum precursor loci per mature miRNA (precursor counts cannot separate the 5p and 3p arms), average replicates and take the median across lines. Gates: every file has at least 1,500 gene rows; miR-21 sits in the panel's top decile (17.8 vs threshold 12.3); 131/131 panel miRNAs map to a precursor. The pre-registered tissue gate failed: we expected miR-122 to peak in HepG2, but it is zero in all 12 lines. We record that failure and do not retune it. A replacement control chosen after the failure (post-hoc, not pre-registered) passes: MIR142's top three lines are the three blood lines (GM12878, HL-60, K562), and MIR302A, MIR302B and MIR367 each peak in H1 embryonic stem cells. All three hypotheses are falsified on $n = 131$ miRNAs: abundance versus CNN AUROC $\rho = 0.076$ ($p = 0.195$), versus the CNN's gain over site-count covariates $\rho = 0.099$ ($p = 0.130$), and the partial correlation given $\log n_{\text{pos}}$ $\rho = 0.050$ ($p = 0.284$; Figure 12). Abundance is also unrelated to the covariate-only AUROC ($\rho = 0.010$). A caveat: ENCORI CLIP pools many cell types, and none of the 12 ENCODE lines is HEK293, so a cell-matched abundance could still matter.

4.21 Tissue target avoidance (GTEx): conserved targets avoid the miRNA's tissue, CNN non-conserved targets do not

Targets of a tissue-specific miRNA should be relatively depleted in that tissue. We test this with GTEx v8 gene median TPM (`scripts/gtex_target_avoidance.py`; `results/gtex_target_avoidance.json`; hermetic tests) on 8 pre-registered pairs (miR-1-3p, miR-133b, miR-206: skeletal muscle; miR-122-5p: liver; miR-9-5p, miR-128-3p: brain cortex; miR-223-3p: whole blood; miR-375: pancreas). Relative




Figure 12: Per-miRNA ENCODE abundance (median over 12 cell lines) versus CNN AUROC on CLIP labels (left) and versus the CNN’s gain over site-count covariates (right).

expression is $r_g(t) = \log_2(\text{TPM}_{g,t} + 1) - \text{median}_{t'} \log_2(\text{TPM}_{g,t'} + 1)$ and depletion is the Mann–Whitney probability $\text{AUC}_{\text{dep}} = P(r_{\text{set}} < r_{\text{rest}})$. Gates passed: ALB peaks in liver, ACTA1 in skeletal muscle, median mapping 95.6%. CNN top-200 genes are not depleted in the matched tissue (H1: 1 wins, 7 losses, $p = 0.996$) and not more depleted there than in the other design tissues (H3: 5/3, $p = 0.363$); both are falsified. Pre-registered H2 (TargetScan genes inside the CLIP universe) turned out to be empty by construction, since that universe excludes every gene with a conserved site for the seed; we record it as not evaluable. A post-hoc reference arm that we did not pre-register uses all TargetScan 3’UTR genes as the universe and the top-200 conserved-site genes as the set: they are depleted in the matched tissue for 8/8 pairs ($p = 0.004$) and more so than in other tissues for 7/8 ($p = 0.035$; Figure 13), which recovers the classic avoidance signal (Farh et al. 2005). So the method sees avoidance where it is expected, and the CNN’s non-conserved candidates show none. Two caveats: the arms use different universes (CNN scores exist only for the CLIP universe), and conservation and scorer are confounded, so this does not separate “non-conserved sites are not functional in vivo” from “the CNN ranks them poorly.” Both readings are negative for the non-conserved discovery claim.




Figure 13: Depletion (AUC_{dep}) of target sets in each miRNA’s own GTEx tissue versus the mean over the other design tissues. Blue: CNN top-200 in the CLIP universe (pre-registered). Red: TargetScan conserved top-200 in the all-UTR universe (post-hoc reference).

4.22 Reverse lookup (Enrichr): CNN top-200 sets name their own miRNA beyond expression, but the site-count win does not replicate

A stricter use of a target list is to ask whether it identifies its miRNA among thousands. We submit each set to the Enrichr API against the miRTarBase_2017 library (3,240 miRNA terms; `scripts/enrichr_reverse_lookup.py`; `results/enrichr_reverse_lookup.json`; hermetic tests) and record the rank of the query miRNA’s own term by Enrichr’s Fisher p -value. Panel, sets and seeds match the g:Profiler section. Gates passed: the library’s own miR-21-5p set returns miR-21-5p at rank 1; the median random-set own-term rank is 283; all 13 miRNAs are library terms. H1 is confirmed: CNN top-200 sets rank their own miRNA better than random for 13/13 miRNAs ($p = 1.2 \times 10^{-4}$; median rank 3), and so do CLIP-positive sets (H2, 13/13; partly circular because miRTarBase includes CLIP-derived entries). H3 is falsified: the HEK293 expression top-200 control also ranks the own term well, and the CNN beats it on only 7/13 (4 losses, $p = 0.274$). Because the CNN’s univariate CLIP signal was earlier shown to be mostly site-count confounding, and expression dominates the other audits, we added two controls after H1 came back positive (post-hoc, not pre-registered). The CNN beats a site-count top-200 set drawn from the same universe (11 wins, 0 losses, 2 ties, $p = 4.9 \times 10^{-4}$) and random sets matched to the CNN set’s expression-decile histogram (13 wins, 0 losses, $p = 1.2 \times 10^{-4}$; Figure 14). Limits: miRTarBase 2017 overlaps the miRTarBase 10.0 labels used in the earlier benchmark, so this is a second view of the same kind of evidence; heavily studied miRNAs (miR-21, miR-155) reach rank 1 even from random sets; and both controls were post-hoc. We therefore pre-registered a held-out replication (`scripts/enrichr_replication.py`, committed before its results; `results/enrichr_replication.json`) on 30 new miRNAs chosen by a seeded shuffle, with both controls fixed in advance. Median own-term ranks were CNN 10.5, site count 70, expression-matched random 433 and uniform random 506. The CNN beats uniform random (26/4, $p = 3.0 \times 10^{-5}$) and expression-matched random (25/5, $p = 1.6 \times 10^{-4}$), but not site count (16 wins, 10 losses, $p = 0.163$). By the pre-registered rule the discovery claim does not replicate: a CNN target list names its miRNA far better than chance and than expression, but most of that comes from seed-site count, as in the per-miRNA feature audit. The site-count win on the discovery panel was a small-panel result.




Figure 14: Rank of each miRNA’s own miRTarBase_2017 term when its target set is submitted to Enrichr plotted as $\log_{10}(3241/\text{rank})$, so taller bars mean better ranks. Site-count and expression-matched arms are post-hoc controls.

4.23 Reactome pathway audit with controls fixed in advance: two gates failed, verdicts inconclusive

We repeated the pathway question with Reactome v97 (AnalysisService projection; `scripts/reactome_enrichment.py`; `results/reactome_enrichment.json`; hermetic tests), this time with the site-count and expression-matched controls pre-registered. Per miRNA we count pathways at entities FDR < 0.05 for CNN top-200, site-count top-200, three expression-matched and three uniform random sets. Two of three gates failed as written. G1 required a pathway named “Cell Cycle” among the top three for a 20-gene cell-cycle control; the control returned 205 significant pathways, but the top three are tied at the same FDR and none has that name (a post-hoc look finds 7 such pathways lower in the list). G3 required fewer than 10% unmapped genes; the median is 39.8%, because Reactome curates only part of the genome. Both are design errors on our side, and we do not rewrite the gates. G2 passed (median uniform-set count 0). Under failed gates we report the numbers but call the audit inconclusive: CNN versus uniform random 3/6 ($p = 0.910$), versus site count 2/2 ($p = 0.688$), versus expression-matched 4/6 ($p = 0.828$). Most sets in every arm have zero significant pathways, with a few large random-set outliers (Figure 15). None of this points toward pathway coherence, which agrees with the g:Profiler result.




Figure 15: Significant Reactome pathways per miRNA for four arms of 200 genes. Gates G1 and G3 failed, so these counts are descriptive.

4.24 Input integrity: every miRNA sequence and seed used here matches RNAcentral

All scorers in this paper key on the miRNA seed (nt 2–8) taken from TargetScan’s `miR_Family_Info`. A wrong sequence would silently corrupt every benchmark, so we checked each miRNA used in the CLIP and miRTarBase analyses against RNAcentral by miRBase accession via the RNAcentral REST API (script and results: `rnacentral_seq_audit`; hermetic tests). Both hypotheses are confirmed (Table 1). One transport note: for a few accessions the endpoint served an HTML page to Python requests but JSON to curl, so the fetcher falls back to curl when the body is not JSON; an earlier run without the fallback left one or two accessions unresolved and was discarded (it was also overwritten by a duplicate process). No analysis choice changed. This rules out sequence errors as an explanation for the negatives above.

Quantity	Value
miRNAs used / with MIMAT accession	132 / 132
Resolved in RNAcentral	132 (100.0%)
G1: let-7a-5p exact sequence	pass
G3: seed identity under deranged map (control)	0.8%
H1: seed identity, resolved miRNAs	100.0%
H2: full-length identity, resolved miRNAs	100.0%

Table 1: RNAcentral sequence audit. The derangement control shows the identity test can fail.

4.25 Independent curated labels (OmniPath, miRTarBase removed): CNN above chance, but not above site count or expression

The miRTarBase and Enrichr analyses share one curation source. OmniPath aggregates several (script and results: `omnipath_independent`; hermetic tests), so we kept only interactions whose sources exclude miRTarBase (2,795 of 11,240; ncRDeathDB, miRecords, miR2Disease, miRDeathDB, SIGNOR) and scored them inside each miRNA’s CLIP universe of non-conserved seed-match genes. The site-count, expression and label-permutation controls were fixed before scoring. Gates passed: 13 miRNAs have at least five positives (bar 10), and the median permutation-mean CNN AUROC is 0.504. The CNN is above chance (H1: 11/2, $p = 0.011$, median AUROC 0.566), but it does not beat site count (H2: 8/5, $p = 0.291$, median 0.572) and loses to HEK293 expression (H3: 3/10, median 0.689; Figure 16). This matches the Enrichr replication on a label set that shares no miRTarBase rows: curated low-throughput targets favour expressed genes, and whatever the CNN adds over counting seed sites is not detectable with 5–29 positives per miRNA.




Figure 16: Per-miRNA AUROC on OmniPath interactions with miRTarBase-sourced rows removed, for the CNN, site count and HEK293 expression.

4.26 Transcription-regulator enrichment (UniProt): present in conserved targets, absent in CNN non-conserved targets

Conserved miRNA targets are known to be enriched for transcription regulators. We test whether CNN non-conserved targets share this property using UniProtKB release 2026_03 (reviewed human entries with keyword KW-0805, 2,374 of 20,190; script and results: `uniprot_tf_enrichment`; hermetic tests). The statistic is the fraction of a 200-gene set carrying KW-0805; site-count and expression-matched controls were fixed in advance. Gates passed. The positive control recovers the classic result: TargetScan conserved

top-200 sets are richer in transcription regulators than random UTR genes for 13/13 miRNAs ($p = 1.2 \times 10^{-4}$), and KW-0805 prevalence is 11.8%. CNN top-200 sets show no enrichment: versus uniform random 6/7 ($p = 0.709$), versus site count 4/9 ($p = 0.954$), versus expression-matched 7/6 ($p = 0.500$); all three are falsified (Figure 17). Together with the GTEx avoidance result this is the second functional signature that conserved targets carry and CNN non-conserved candidates lack. The same caveat applies: conservation and scorer are confounded, and the universes differ between arms.




Figure 17: Transcription-regulator fraction of each 200-gene set minus its random baseline (percentage points). Red: TargetScan conserved targets against random UTR genes (positive control).

4.27 Loss-of-function constraint (gnomAD): conserved targets are constrained, CNN candidates are the least constrained set

Conserved miRNA targets tend to be dosage-sensitive. We compared the median LOEUF (gnomAD v2.1.1 upper bound of observed/expected loss-of-function; lower means more constrained; 19,155 genes; script and results: `gnomad_constraint`) of 200-gene sets, with the same controls fixed in advance. Gates passed: 97.6% of CLIP-universe genes have a LOEUF, and TargetScan conserved top-200 sets are more constrained than random UTR genes for 13/13 miRNAs ($p = 1.2 \times 10^{-4}$). All three pre-registered CNN hypotheses are falsified, and in the opposite direction: the CNN top-200 has a higher median LOEUF than uniform random (13/13), site count (13/13) and expression-matched sets (13/13; Figure 18). Not pre-registered: the two-sided sign test of this reversal gives $p = 2.4 \times 10^{-4}$ against uniform and $p = 2.4 \times 10^{-4}$ against expression-matched sets. Site-count sets are more constrained than uniform random for 10/13 miRNAs, as expected if they favour genes with long 3' UTRs. The CNN's preference for genes with few, strong sites may select short-UTR, less constrained genes. This post-hoc finding needs a held-out, UTR-length-matched replication before it can be read as biology, and we make no claim beyond the description.

Pre-registered held-out replication. We committed a replication script before running it (`gnomad_replication`): 30 CLIP miRNAs outside the discovery panel, a one-sided test in the discovered direction, and a new UTR-length-matched control (decile matching on the longest 3' UTR). Gates passed (97.8% coverage). The reversal replicates: the CNN top-200 is less constrained than uniform random in 30/30 miRNAs ($p = 9.3 \times 10^{-10}$; median LOEUF difference +0.192), than expression-matched sets in 30/30 ($p = 9.3 \times 10^{-10}$; +0.161), and than UTR-length-matched sets in 25/30 ($p = 1.6 \times 10^{-4}$; +0.084; Figure 19). CNN top-200 genes have shorter UTRs than their universe in 30/30 miRNAs, and length matching shrinks the gap, but it does not




Figure 18: Median LOEUF of each 200-gene set. Black dashes: random UTR genes, the comparator for the TargetScan positive control.

remove it. At that point we treated this as a falsifiable candidate pattern, not a confirmed mechanism. Subsequent pre-registered dosage-label tests (DECIPHER, ClinGen, HPO) did not support the clause about UTR length, so we withdrew it (Table 9). The later locked random-weight falsifier also failed: untrained weights reproduce the apparent depletion at the required panel-count threshold. We therefore withdraw the learned-model biological claim entirely, despite this descriptive held-out-gene pattern. A finer absolute UTR-length-bin audit, run after the failed falsifier, is reported separately below; it cannot restore the failed gate. Purifying selection was only an untested hypothesis, not an inference here.




Figure 19: Pre-registered replication on 30 held-out miRNAs. Left: CNN top-200 median LOEUF minus each control (positive = CNN less constrained). Right: CNN top-200 median UTR length relative to its universe.

4.28 Independent dosage label (ClinGen): the depletion beyond UTR length does not replicate

gnomAD LOEUF is a population statistic, so we tested the named finding against an independent label: expert-curated ClinGen haploinsufficient genes (score 3, 420 genes), on all 43 miRNAs (13 discovery + 30 held-out). The script was committed before it was run (`clingen_dosage`). Gates passed: the pooled CLIP universe holds 394 HI genes, and TargetScan conserved top-200 sets hold more HI genes than random

UTR genes (32/5 miRNAs, $p = 3.7 \times 10^{-6}$; pooled 307 vs 166). The CNN sets hold fewer HI genes than

Pre-registered test	wins/losses	p	pooled CNN / control	verdict
C1 CNN < uniform	26/16	0.082	158 / 205.3	FALSIFIED
C2 CNN < expression-matched	28/14	0.022	158 / 191.7	CONFIRMED
C3 CNN < UTR-length-matched	24/16	0.134	158 / 159.3	FALSIFIED

Table 2: ClinGen HI genes in CNN top-200 sets. Per-miRNA sign tests (ties excluded) and pooled counts over 43 miRNAs.

expression-matched sets (C2 confirmed; pooled ratio 0.82), but the depletion against uniform sets is not significant (C1) and it vanishes against UTR-length-matched sets (pooled ratio 0.99; C3 falsified; Figure 20). On this independent label, UTR length explains the depletion. The gnomAD finding therefore survives only in its weaker form (CNN candidates are less constrained than expression-matched genes); its “beyond UTR length” clause is supported by LOEUF but not by ClinGen. HI genes are sparse (about 4 per set), so C3 has little power; we report the disagreement rather than choosing between the labels.




Figure 20: Pooled HI-gene count in each set divided by its control. Red: positive control.

4.29 Third dosage label (Collins et al. 2022 rCNV scores): the depletion survives UTR-length matching

We then used the pHaplo and pTriplo scores of Collins et al. (Zenodo 6347673; 18,641 genes), which are estimated from rare CNVs in about 950,000 people. We noted in advance that the model uses gene features that include gnomAD constraint, so this label is only partly independent of LOEUF. Same 43 miRNAs and controls; the script was committed before it was run (collins_dosage). Gates passed: 95.4% coverage, and TargetScan conserved top-200 sets have higher pHaplo than random UTR genes (36/2, $p = 2.7 \times 10^{-9}$). All four hypotheses are confirmed, including against UTR-length-matched sets for both haploinsufficiency

Pre-registered test (CNN lower)	wins/losses	p	median diff.	verdict
D1 pHaplo vs uniform	43/0	1.1×10^{-13}	-0.131	CONFIRMED
D2 pHaplo vs expression-matched	43/0	1.1×10^{-13}	-0.099	CONFIRMED
D3 pHaplo vs UTR-length-matched	38/5	1.2×10^{-7}	-0.038	CONFIRMED
T3 pTriplo vs UTR-length-matched	42/1	5.0×10^{-12}	-0.054	CONFIRMED

Table 3: Collins 2022 dosage scores of CNN top-200 sets, 43 miRNAs.

and triplosensitivity (Figure 21). Across the three labels the picture is: LOEUF and the rCNV scores

support depletion beyond UTR length; the sparse ClinGen curation does not. Because the rCNV model shares features with LOEUF, it is not a clean third vote, and the finding stays at the strength of its weakest independent test. The later locked random-weight falsifier failed, so these descriptive dosage-label contrasts are not evidence that trained CNN weights caused the pattern.




Figure 21: Per-miRNA difference between CNN top-200 and control median dosage score (negative = CNN less dosage-sensitive).

4.30 Mendelian disease genes (HPO): depletion explained by UTR length

HPO gene–disease annotations give a fourth, curated label (4,886 genes with a Mendelian association; `hpo_disease_genes`, committed before running; same 43 miRNAs and controls). Gates passed: disease-gene prevalence in the pooled universe is 24.1%, and TargetScan conserved top-200 sets hold more disease genes than random UTR genes (27/11, $p = 0.007$). CNN sets carry fewer disease genes than uniform and

Pre-registered test	wins/losses	p	pooled CNN / control (ratio)	verdict
M1 CNN fewer than uniform	30/10	0.001	1971 / 2175 (0.91)	CONFIRMED
M2 fewer than expression-matched	27/15	0.044	1971 / 2070 (0.95)	CONFIRMED
M3 fewer than UTR-length-matched	22/21	0.500	1971 / 2075 (0.95)	FALSIFIED

Table 4: HPO Mendelian disease genes in CNN top-200 sets, 43 miRNAs.

expression-matched sets, but against UTR-length-matched sets the per-miRNA split is 22/21 (M3 falsified), even though the pooled ratio stays at 0.95. This agrees with ClinGen: both curated labels support depletion relative to expression, and neither supports it beyond UTR length. The two population-derived scores (LOEUF, rCNV) do. We state the finding in its supported form: *CNN top-ranked non-conserved candidates are depleted of dosage-sensitive and disease genes relative to expression-matched genes*; whether this goes beyond UTR length depends on the label and remains open.

4.31 ClinVar pathogenic-variant genes: the positive-control gate failed

The last label was ClinVar (gene summary dated September 23, 2026): genes with at least five Pathogenic/Likely pathogenic alleles (5,205 genes; `clinvar_pathogenic`, committed before running). We noted in advance that allele counts grow with coding length and testing effort. Prevalence passed (21.8%), but the positive-control gate failed: TargetScan conserved top-200 sets do not hold significantly more PLP genes than random UTR genes (21/15, $p = 0.203$). Because the label does not recover the known conserved-target signal,

Pre-registered test (gate G1 failed)	wins/losses	p	pooled ratio	verdict
V1 fewer than uniform	32/11	9.6×10^{-4}	0.91	CONFIRMED
V2 fewer than expression-matched	29/13	0.010	0.92	CONFIRMED
V3 fewer than UTR-length-matched	24/18	0.220	0.97	FALSIFIED
S3 (≥ 1 PLP) vs UTR-length-matched	23/19	0.322	0.98	FALSIFIED

Table 5: ClinVar PLP genes in CNN top-200 sets, 43 miRNAs. Reported for completeness only.

we treat the CNN results below as uninterpretable and do not count them for or against the finding. For completeness, the pattern matches ClinGen and HPO (depletion against uniform and expression-matched sets, none against UTR-length-matched sets).

4.32 Cell-essential genes (DepMap): no depletion, and the label is uninformative for conserved targets

To ask whether the dosage depletion extends to cell fitness, we used the DepMap 24Q4 inferred common-essential genes (1,523 genes; `depmap_essential`, committed before running; same 43 miRNAs and controls). Prevalence passed (7.3%). The positive-control gate, which only asked for a difference in either direction, failed: conserved top-200 sets and random UTR genes split 20/17 (two-sided $p = 0.743$). All

Pre-registered test	wins/losses	p	pooled CNN / control (ratio)	verdict
E1 fewer than uniform	19/21	0.682	662 / 634 (1.04)	FALSIFIED
E2 fewer than expression-matched	14/29	0.993	662 / 603 (1.10)	FALSIFIED
E3 fewer than UTR-length-matched	18/24	0.860	662 / 627 (1.06)	FALSIFIED

Table 6: DepMap common-essential genes in CNN top-200 sets, 43 miRNAs.

three hypotheses are falsified. CNN sets hold slightly more essential genes than every control, not fewer (pooled ratios 1.04–1.10). Whatever drives the depletion of dosage-sensitive and disease genes, it does not extend to genes needed for cell growth in culture. That fits a dosage reading better than a general avoidance of important genes, but with the positive-control gate failed we do not lean on it.

4.33 Protein-complex subunits (Complex Portal): the dosage-balance prediction fails

The dosage-balance hypothesis holds that subunits of protein complexes are dosage-sensitive, so the dosage finding predicts that CNN sets carry fewer subunit genes. We used EBI Complex Portal (2,498 human complexes, 3,686 subunit genes after mapping UniProt accessions; `complexportal_subunits`, committed before running; same 43 miRNAs and controls). Gates passed: 99.9% of accessions mapped, prevalence 17.9%, and conserved top-200 sets differ from random UTR genes (27/9 more subunits, two-sided $p = 0.004$). The depletion against uniform sets disappears once expression or UTR length is matched (X2, X3

Pre-registered test	wins/losses	p	pooled CNN / control (ratio)	verdict
X1 fewer than uniform	29/12	0.006	1497 / 1572 (0.95)	CONFIRMED
X2 fewer than expression-matched	19/23	0.780	1497 / 1467 (1.02)	FALSIFIED
X3 fewer than UTR-length-matched	19/24	0.820	1497 / 1503 (1.00)	FALSIFIED

Table 7: Complex Portal subunit genes in CNN top-200 sets, 43 miRNAs.

falsified; pooled ratios 1.02 and 1.00). So the CNN does not avoid complex subunits beyond what expression explains, and the dosage-balance route does not account for the LOEUF and rCNV depletion.

4.34 Study bias (NCBI Gene, PubMed counts): it explains the curated-label depletion but not the LOEUF depletion

Curated labels (ClinGen, HPO, ClinVar) favour well-studied genes. We counted PubMed records per human gene from NCBI Gene (gene_info and gene2pubmed; ncbi_pubmed_bias, committed before running) and added a publication-matched control (deciles of $\log_{10}(1+n)$). The positive control passed (conserved targets are better studied, 37/1, $p = 1.4 \times 10^{-10}$), but the mapping gate failed: 89.8% of universe symbols matched an NCBI symbol against a 95% bar. The pre-registered script scores unmapped symbols as zero publications, so we add a sensitivity analysis, labelled not pre-registered, that drops them (ncbi_pubmed_posthoc). Unmapped symbols are more common in CNN sets (35/8 miRNAs; median 11.0% vs 9.2%). Both analyses

Test	wins/losses	p	verdict (G2 failed)	unmapped dropped [†]
P1 CNN less studied than uniform	40/3	1.5×10^{-9}	CONFIRMED	38/5 (1.2×10^{-7})
P2 less studied than expression-matched	35/8	2.1×10^{-5}	CONFIRMED	38/5 (1.2×10^{-7})
P3 less studied than UTR-length-matched	36/7	4.5×10^{-6}	CONFIRMED	33/10 (3.0×10^{-4})
P4 LOEUF higher than publication-matched	43/0	1.1×10^{-13}	CONFIRMED	43/0 (1.1×10^{-13})
P5 fewer HPO genes than publication-matched	20/23	0.729	FALSIFIED	15/27 (0.978)

Table 8: Study bias of CNN top-200 sets, 43 miRNAs. [†]Not pre-registered.

agree. CNN candidates are less studied than every control. Against publication-matched sets the LOEUF depletion stays at 43/0, so study bias does not explain it, which is expected because LOEUF comes from population sequencing, not from the literature. The HPO depletion disappears (P5 falsified, and reversed in the post-hoc split). This explains why the curated labels (ClinGen, HPO) disagreed with LOEUF: their depletion signal is largely study bias. Because gate G2 failed, the verdicts carry that caveat; the post-hoc split is not a substitute for a pre-registered test.

4.35 DECIPHER haploinsufficiency predictions: depletion beyond expression and study bias, not beyond UTR length

DECIPHER’s haploinsufficiency predictions (version 3; 19,124 genes) come from a model of genomic, evolutionary and network features that does not use gnomAD LOEUF (decipher_hi, committed before running; 43 miRNAs; publication-matched control added). Gates passed: 95.9% coverage, and conserved top-200 sets score higher than random UTR genes (38/0, $p = 3.6 \times 10^{-12}$). The CNN depletion holds against

Pre-registered test	wins/losses	p	median diff.	verdict
H1 lower than uniform	42/1	5.0×10^{-12}	-0.0554	CONFIRMED
H2 lower than expression-matched	41/2	1.1×10^{-10}	-0.0448	CONFIRMED
H3 lower than UTR-length-matched	25/18	0.180	-0.0039	FALSIFIED
H4 lower than publication-matched	39/4	1.6×10^{-8}	-0.0400	CONFIRMED

Table 9: DECIPHER HI scores of CNN top-200 sets, 43 miRNAs.

expression-matched and publication-matched sets but not against UTR-length-matched sets (H3 falsified; median difference -0.0039). The tally for “beyond UTR length” is now: supported by LOEUF and the rCNV scores, not supported by DECIPHER, ClinGen or HPO (ClinVar and DepMap uninformative). We therefore do not claim the depletion goes beyond UTR length. The claim that survives every informative label is the weaker one: CNN top-ranked non-conserved candidates are less dosage-sensitive than expression-matched genes.

4.36 GeneBayes s_{het} : population-based labels agree with each other, curated and model-based labels do not

GeneBayes s_{het} (Zenodo 10403680; 19,068 genes) estimates selection against heterozygous loss of function from gnomAD LoF counts plus a gene-feature prior, so it is a second population-based label and not independent of LOEUF (genebayes_shet, committed before running; statistic median $\log_{10} s_{\text{het}}$). Gates passed: 96.8% coverage, and the positive control scores higher than random (38/0, $p = 3.6 \times 10^{-12}$). All

Pre-registered test	wins/losses	p	median diff.	verdict
H1 lower than uniform	43/0	1.1×10^{-13}	-0.416	CONFIRMED
H2 lower than expression-matched	43/0	1.1×10^{-13}	-0.389	CONFIRMED
H3 lower than UTR-length-matched	39/4	1.6×10^{-8}	-0.169	CONFIRMED
H4 lower than publication-matched	43/0	1.1×10^{-13}	-0.367	CONFIRMED

Table 10: GeneBayes s_{het} of CNN top-200 sets, 43 miRNAs (differences in $\log_{10}$ units).

four hypotheses are confirmed, including H3 against UTR-length-matched sets. The split is now clean: the three labels built from gnomAD loss-of-function counts (LOEUF, rCNV, s_{het}) support depletion beyond UTR length, while DECIPHER, ClinGen and HPO do not. A plausible explanation, not tested here, is that labels built from LoF counts are less informative for genes with few expected LoF variants (short coding sequences), and the CNN may favour such genes. Because the split follows label construction, the clause about UTR length stays withdrawn; the surviving claim (depletion relative to expression-matched and publication-matched genes) is confirmed here too.

Pre-registered coding-length control. We tested the explanation above directly (explof_matched, committed before running), matching on gnomAD’s expected LoF count, a proxy for coding length. CNN top-200 genes do have fewer expected LoF variants than their universe (41/2 miRNAs), but only by a median factor of 0.92. The matching worked (matched/CNN expected-LoF ratio 0.999). The depletion survives: CNN LOEUF is higher than in expected-LoF-matched sets for 43/0 miRNAs ($p = 1.1 \times 10^{-13}$) and s_{het} lower for 42/1 ($p = 5.0 \times 10^{-12}$). So the coding-length explanation is falsified. The disagreement between LoF-count labels and the other labels stays unexplained, and the clause about UTR length stays withdrawn.

4.37 Paralog buffering (Ensembl BioMart): no paralog excess; both gates failed

Genes with paralogs can be buffered against dosage changes, so one explanation of the depletion is that CNN candidates have more paralogs. We counted human paralogues per gene from Ensembl BioMart (biomart_paralogs, committed before running; 43 miRNAs; statistic median $\log_2(1 + \text{paralogs})$). Both gates failed: 89.5% of universe symbols were found (bar 90%), and conserved targets did not differ from random genes (12/24, two-sided $p = 0.065$). CNN candidates have no paralog excess (B1, B2 falsified),

Pre-registered test (G1, G2 failed)	wins/losses	p	median diff.	verdict
B1 more paralogs than uniform	22/20	0.439	+0.014	FALSIFIED
B2 more paralogs than expression-matched	17/24	0.894	-0.044	FALSIFIED
B3 LOEUF higher than paralog-matched	43/0	1.1×10^{-13}	+0.187	CONFIRMED

Table 11: Paralog counts of CNN top-200 sets, 43 miRNAs.

and the LOEUF depletion stays at 43/0 against paralog-matched sets. With both gates failed we give these results little weight, but nothing here supports paralog buffering as the explanation.

4.38 Gene families (HGNC): the depletion is not a zinc-finger artefact; coverage gate failed

Large, dosage-tolerant families such as the C2H2 zinc fingers (many with long 3'UTRs) could make a candidate set look unconstrained. We used HGNC gene groups [Seal et al.(2023)] (19297 approved protein-coding symbols, 763 C2H2-ZNF; hgnc_families, committed before running; 43 miRNAs; family size = $\log_2(1 + \text{largest group size})$). The parse gate passed (C2H2-ZNF genes are $7.1 \times$ more frequent on chr19, bar $3 \times$), but the coverage gate failed: 88.7% of universe symbols matched current HGNC protein-coding symbols (bar 90%). A post-hoc audit (not pre-registered) found that 1311 of the 1974 unmatched symbols are previous HGNC symbols and 195 are approved non-protein-coding loci. CNN candidates are not

Pre-registered test (G1 failed)	wins/losses	p	median diff.	verdict
F1 C2H2-ZNF fraction above expression-matched	21/22	0.620	+0.000	FALSIFIED
F2 family size above expression-matched	21/22	0.620	-0.025	FALSIFIED
F3 LOEUF above expression-matched, ZNFs removed	43/0	1.1×10^{-13}	+0.152	CONFIRMED

Table 12: HGNC family composition of CNN top-200 sets, 43 miRNAs.

enriched for zinc fingers or large families (F1, F2 falsified), and with every C2H2-ZNF gene removed the LOEUF depletion against expression-matched sets stays at 43/0. Because the coverage gate failed, this rules out the family explanation only for the matched genes.

4.39 Orphanet rare-disease genes: depletion vs expression-matched only; gone against publication-matched sets

Orphanet curates disorder-gene links by hand and records the mechanism, so it gives a further disease label and an explicit loss-of-function label [Orphanet(2024)]. From Orphadata en_product6 we took assessed germline disease-causing links (4195 genes) and the loss-of-function subset (1047 genes); statistic = fraction of the top-200 (orphanet_genes, committed before running; 43 miRNAs). Both gates passed: conserved TargetScan targets carry more disease genes than random UTR genes (27/11, $p = 0.007$), and the parse counts met the bars. CNN candidates carry fewer disease genes than expression-matched sets (O1), but

Pre-registered test (CNN lower)	wins/losses	p	median diff.	verdict
O1 disease genes, vs expression-matched	30/13	0.007	-0.018	CONFIRMED
O2 disease genes, vs publication-matched	21/21	0.561	+0.000	FALSIFIED
O3 LoF genes, vs publication-matched	20/23	0.729	+0.000	FALSIFIED
O4 LoF genes, vs UTR-length-matched	20/23	0.729	+0.002	FALSIFIED

Table 13: Orphanet disease-gene fractions in CNN top-200 sets, 43 miRNAs.

the gap disappears against publication-matched sets (O2), and the explicit loss-of-function label shows no depletion against publication- or UTR-length-matched sets (O3, O4). This repeats the HPO pattern: on curated labels the depletion is explained by study bias, while the LoF-count labels (LOEUF, s_{het}) keep it.

4.40 Mouse mutant phenotypes (MGI): weak, mixed support

Mouse knockouts give an experimental label that owes nothing to human population data or clinical curation. From the MGI HMD.HumanPhenotype report [Baldarelli et al.(2024)] we marked human genes whose mouse orthologue mutants show mortality/aging (MP:0010768) or embryo (MP:0005380) phenotypes (6603 and 2429 of 14268 phenotyped genes); genes never phenotyped were dropped (mgi_mouse_ko, committed before running; 43 miRNAs). Both gates passed: conserved TargetScan targets are more often lethal than

random genes (37/1, $p = 1.4 \times 10^{-10}$), and 85.4% of universe symbols have an MGI row (bar 80%). All

Pre-registered test (CNN lower)	wins/losses	p	median diff.	verdict
M1 mortality, vs expression-matched	28/14	0.022	-0.013	CONFIRMED
M2 mortality, vs publication-matched	26/16	0.082	-0.022	FALSIFIED
M3 mortality, vs UTR-length-matched	26/17	0.111	-0.013	FALSIFIED
M4 embryo phenotype, vs publication-matched	29/14	0.016	-0.012	CONFIRMED

Table 14: Mouse mutant phenotype fractions in CNN top-200 sets, 43 miRNAs.

four comparisons point the same way (CNN sets less often lethal), but only M1 and M4 reach $p < 0.05$; mortality is not significantly depleted against publication- or UTR-length-matched sets. The plan set no multiple-testing correction; after the fact we note that with Bonferroni over the four tests ($\alpha = 0.0125$) none would pass. We read this as weak support from an independent experimental label, not as a replication.

5 Independent benchmark audit: every evaluated miRNA

The independent miRTarBase 10.0 label benchmark is the clearest competitive boundary for this method. On 7,909 matched rows, miRDB v6.0 has pooled AUROC 0.7867 versus DuplexCNN 0.6847; miRDB wins on 74 of 84 miRNAs. The tables below preserve each per-miRNA sample count, positive count, and pair of AUROCs from `results/mirdb_benchmark.json`. These are existing observations, not 84 independent studies or newly run experiments. The smaller miRNAs have unstable estimates. Winning individual rows cannot reverse the pooled or paired overall result.

miRNA	Pairs	Positive	CNN AUROC	miRDB AUROC	Δ (CNN-miRDB)
hsa-let-7a-5p	86	24	0.577	0.833	-0.257
hsa-miR-1-3p	176	49	0.746	0.898	-0.152
hsa-miR-100-5p	14	9	0.822	0.833	-0.011
hsa-miR-103a-3p	34	9	0.707	0.707	+0.000
hsa-miR-106a-5p	103	25	0.610	0.772	-0.162
hsa-miR-106b-5p	115	24	0.633	0.889	-0.255
hsa-miR-107	70	20	0.702	0.782	-0.080
hsa-miR-10b-5p	52	15	0.802	0.751	+0.050
hsa-miR-122-5p	175	43	0.720	0.658	+0.061
hsa-miR-125b-5p	236	58	0.663	0.775	-0.112
hsa-miR-128-3p	89	18	0.703	0.897	-0.194
hsa-miR-129-5p	47	9	0.482	0.693	-0.211
hsa-miR-130a-3p	56	16	0.737	0.788	-0.050
hsa-miR-130b-3p	39	11	0.808	0.899	-0.091
hsa-miR-132-3p	28	10	0.750	0.800	-0.050
hsa-miR-133b	58	18	0.664	0.825	-0.161
hsa-miR-135b-5p	29	9	0.667	0.789	-0.122
hsa-miR-138-5p	82	23	0.730	0.809	-0.079
hsa-miR-140-5p	33	7	0.783	0.989	-0.206
hsa-miR-141-3p	148	34	0.555	0.721	-0.166
hsa-miR-143-3p	83	20	0.621	0.743	-0.122

Table 15: Independent miRTarBase label audit, miRNAs 1–21 of 84, in source-record order. Scores are unweighted per-miRNA AUROC and not a pooled estimate.

5.1 Interpretation of the independent ranking

The source data include 2,113 positive rows out of 7,909, with miRDB scores present for 29.8% of rows. The originally declared test labels, missing-score handling and benchmark construction determine the meaning

miRNA	Pairs	Positive	CNN AUROC	miRDB AUROC	Δ (CNN–miRDB)
hsa-miR-144-3p	48	11	0.740	0.765	-0.026
hsa-miR-145-5p	345	79	0.683	0.748	-0.065
hsa-miR-146a-5p	76	16	0.658	0.660	-0.002
hsa-miR-148a-3p	77	22	0.707	0.819	-0.112
hsa-miR-150-5p	81	14	0.643	0.834	-0.191
hsa-miR-152-3p	33	9	0.708	0.884	-0.176
hsa-miR-155-5p	419	167	0.770	0.761	+0.009
hsa-miR-15a-5p	102	29	0.778	0.870	-0.092
hsa-miR-15b-5p	52	12	0.729	0.827	-0.098
hsa-miR-16-5p	148	36	0.677	0.867	-0.191
hsa-miR-17-5p	220	49	0.649	0.893	-0.244
hsa-miR-181a-5p	158	36	0.651	0.833	-0.183
hsa-miR-182-5p	116	29	0.731	0.791	-0.060
hsa-miR-185-5p	74	13	0.657	0.788	-0.131
hsa-miR-18a-5p	25	7	0.571	0.877	-0.306
hsa-miR-191-5p	20	12	0.750	0.792	-0.042
hsa-miR-192-5p	60	19	0.675	0.575	+0.100
hsa-miR-194-5p	44	15	0.676	0.736	-0.060
hsa-miR-195-5p	95	26	0.773	0.829	-0.056
hsa-miR-196a-5p	30	12	0.792	0.861	-0.069
hsa-miR-199a-5p	83	20	0.719	0.736	-0.017

Table 16: Independent miRTarBase label audit, miRNAs 22–42 of 84, in source-record order. Scores are unweighted per-miRNA AUROC and not a pooled estimate.

miRNA	Pairs	Positive	CNN AUROC	miRDB AUROC	Δ (CNN–miRDB)
hsa-miR-19a-3p	91	27	0.716	0.897	-0.181
hsa-miR-19b-3p	61	18	0.583	0.848	-0.266
hsa-miR-200a-3p	109	21	0.575	0.721	-0.146
hsa-miR-200b-3p	137	37	0.625	0.652	-0.027
hsa-miR-200c-3p	178	46	0.712	0.797	-0.085
hsa-miR-204-5p	153	36	0.756	0.779	-0.023
hsa-miR-205-5p	108	26	0.744	0.821	-0.077
hsa-miR-206	37	15	0.594	0.836	-0.242
hsa-miR-20a-5p	181	38	0.662	0.849	-0.187
hsa-miR-21-5p	162	54	0.693	0.742	-0.049
hsa-miR-210-3p	47	26	0.603	0.696	-0.093
hsa-miR-214-3p	128	25	0.662	0.819	-0.157
hsa-miR-218-5p	97	28	0.724	0.790	-0.066
hsa-miR-22-3p	75	21	0.738	0.809	-0.071
hsa-miR-221-3p	105	34	0.746	0.726	+0.020
hsa-miR-222-3p	68	19	0.802	0.760	+0.042
hsa-miR-223-3p	71	25	0.658	0.841	-0.183
hsa-miR-224-5p	58	17	0.690	0.537	+0.153
hsa-miR-23a-3p	91	27	0.748	0.837	-0.090
hsa-miR-24-3p	184	31	0.624	0.743	-0.119
hsa-miR-25-3p	41	14	0.833	0.770	+0.063

Table 17: Independent miRTarBase label audit, miRNAs 43–63 of 84, in source-record order. Scores are unweighted per-miRNA AUROC and not a pooled estimate.

miRNA	Pairs	Positive	CNN AUROC	miRDB AUROC	Δ (CNN-miRDB)
hsa-miR-26a-5p	107	30	0.721	0.829	-0.107
hsa-miR-27a-3p	90	23	0.734	0.786	-0.052
hsa-miR-28-5p	29	10	0.600	0.784	-0.184
hsa-miR-29a-3p	182	53	0.731	0.874	-0.143
hsa-miR-302a-3p	54	9	0.793	0.804	-0.011
hsa-miR-30a-5p	84	21	0.668	0.815	-0.147
hsa-miR-31-5p	119	32	0.756	0.764	-0.009
hsa-miR-335-5p	59	16	0.590	0.703	-0.113
hsa-miR-34a-5p	203	50	0.768	0.855	-0.087
hsa-miR-378a-3p	36	8	0.638	0.576	+0.062
hsa-miR-424-5p	103	25	0.834	0.927	-0.093
hsa-miR-429	54	16	0.590	0.673	-0.082
hsa-miR-494-3p	65	18	0.577	0.683	-0.106
hsa-miR-503-5p	48	18	0.637	0.708	-0.071
hsa-miR-7-5p	69	19	0.693	0.763	-0.071
hsa-miR-9-5p	126	31	0.571	0.796	-0.224
hsa-miR-92a-3p	81	18	0.675	0.783	-0.108
hsa-miR-93-5p	64	14	0.754	0.924	-0.170
hsa-miR-96-5p	49	14	0.673	0.894	-0.220
hsa-miR-98-5p	30	10	0.610	0.725	-0.115
hsa-miR-99a-5p	16	9	0.683	0.873	-0.190

Table 18: Independent miRTarBase label audit, miRNAs 64–84 of 84, in source-record order. Scores are unweighted per-miRNA AUROC and not a pooled estimate.

of the comparison; the per-miRNA list must not be filtered after seeing scores. The Wilcoxon paired $p = 2.45 \times 10^{-12}$ supports the direction of miRDB advantage across the tested miRNAs, not biological validation of individual target pairs. Later transfer and random-weight controls are separate questions; they cannot rehabilitate a false state-of-the-art claim on this task.

6 Documented negatives

(1) Sequence-only DuplexCNN does not beat ridge. (2) CLIP effect size is small. (3) miR-125b-5p and miR-7-5p per-miRNA AUROCs (negated convention) are at/below chance — reported per miRNA, not averaged away. (4) Pair-map branch contributes only $\Delta r = +0.0029$ (ablation). (5) miRDB v6.0 beats our CNN on miRTarBase rows (pooled AUROC 0.787 vs 0.685, 74/84 miRNAs, $p = 2.4 \times 10^{-12}$). (6) On TarBase v9.0 labels the CNN loses to miRDB on 73/83 miRNAs and to the plain seed-site count per miRNA (CNN better on only 32/83, $p = 0.0078$). (7) HEK293 expression alone (HPA nTPM, AUROC 0.682) beats all sequence models on CLIP labels and cuts the CNN’s conditional pooled gain from +0.0048 to +0.0008. (8) g:Profiler finds no pathway coherence in CNN top-200 (13/13 miRNAs zero terms) or CLIP top-200 target sets; an expression top-200 control gives 47–115 terms. (9) STRING network density: CNN top-200 no denser than random; the CLIP+ excess over random (13/13) does not survive an expression-matched control (8 wins, 4 losses, $p = 0.19$). (10) ENCODE miRNA abundance (12 lines) does not predict per-miRNA CNN AUROC ($p = 0.076$, $p = 0.195$) or gain; pre-registered miR-122/HepG2 gate failed (zero in all lines). (11) GTEx: CNN non-conserved top-200 targets show no tissue avoidance (1 win, 7 losses) and no tissue specificity (5/3, $p = 0.36$); pre-registered TargetScan arm empty by construction. (12) Enrichr: HEK293 expression top-200 names the right miRNA nearly as well as the CNN (H3 falsified, 7/4); the CNN’s reverse-lookup edge over site count (11/0 on the discovery panel) failed pre-registered held-out replication (16/10, $p = 0.16$, 30 miRNAs). (13) Reactome audit: pre-registered gates G1 (tie-ordered top-3 rule) and G3 (40% unmapped) failed; verdicts inconclusive, no sign of pathway coherence. (14) OmniPath non-miRTarBase labels: CNN above chance (11/2) but not above site count (8/5, $p = 0.29$) and below HEK293 expression (3/10). (15) UniProt KW-0805: conserved targets enriched for transcription regulators (13/0), CNN

non-conserved targets not (6/7 vs random, 4/9 vs site count). (16) gnomAD LOEUF: conserved targets constrained (13/0); CNN top-200 less constrained than every control (0/13). The reversal was found post hoc; its pre-registered held-out replication (30 miRNAs) confirmed it against uniform (30/0), expression-matched (30/0) and UTR-length-matched (25/5) sets. A later post-hoc fixed-width UTR-length-bin audit is reported below, but the random-weight falsifier failed and the trained-model claim is withdrawn. (17) ClinGen HI genes (independent label, pre-registered, 43 miRNAs): the CNN depletion holds against expression-matched sets (28/14) but not against uniform (26/16, $p = 0.082$) or UTR-length-matched sets (pooled ratio 0.99). The gnomAD clause *beyond UTR length* is not supported by ClinGen. (18) HPO Mendelian disease genes (pre-registered, 43 miRNAs): CNN depletion holds against uniform (30/10) and expression-matched sets (27/15) but not UTR-length-matched sets (22/21). (19) ClinVar PLP genes: pre-registered positive-control gate failed (21/15, $p = 0.203$); CNN results treated as uninterpretable. (20) DepMap common essentials: no CNN depletion (E1–E3 falsified; pooled ratios above 1) and the positive-control gate failed (20/17). (21) Complex Portal subunits: CNN depletion against uniform (29/12) vanishes against expression-matched (19/23) and UTR-length-matched sets (19/24); the dosage-balance prediction fails. (22) NCBI PubMed counts: pre-registered mapping gate failed (89.8% vs 95%); CNN candidates are less studied than all controls, and the HPO depletion disappears against publication-matched sets (20/23), so the curated-label signal is largely study bias; the LOEUF depletion survives (43/0). (23) DECIPHER HI predictions: CNN depletion holds against expression-matched (41/2) and publication-matched sets (39/4) but not UTR-length-matched sets (25/18, $p = 0.180$); the “beyond UTR length” clause is withdrawn. (24) Coding-length explanation of the LoF-count-label split (pre-registered): falsified; the LOEUF and s_{het} depletion survives expected-LoF matching (43/0, 42/1), and the split between labels remains unexplained. (25) BioMart paralogs: both pre-registered gates failed (coverage 89.5%, positive control $p = 0.065$); no paralog excess in CNN sets (22/20, 17/24). (26) HGNC gene families: pre-registered coverage gate failed (88.7% vs 90%); CNN candidates are not enriched for C2H2 zinc fingers or large families (21/22, 21/22), and the LOEUF depletion survives removal of all C2H2-ZNF genes (43/0). (27) Orphanet disease genes (pre-registered, gates passed): CNN depletion against expression-matched sets (30/13) disappears against publication-matched sets (21/21); no depletion of Orphanet loss-of-function genes against publication- (20/23) or UTR-length-matched sets (20/23). (28) MGI mouse mutant phenotypes (pre-registered, gates passed): mortality depletion holds against expression-matched sets (28/14) but not publication- (26/16, $p = 0.082$) or UTR-length-matched sets (26/17, $p = 0.111$); embryo-phenotype depletion against publication-matched sets (29/14, $p = 0.016$) would not survive a Bonferroni correction.

6.1 Locked random-weight and score-permutation falsifier

The dosage-depletion interpretation faced an independently locked test, separate from the primary expression-controlled CLIP endpoint. We used 200,000 seed-0-subsampled non-conserved candidates, 43-panel miRNAs, the top 200 per miRNA, and three expected-LoF-matched comparison draws. The trained network showed depletion in 31 panels by LOEUF (L1) and 33 by s_{het} (L2). Under 100 independently initialized, *untrained* DuplexCNNs, the 99th-percentile panel counts were 33 for L1 and 33 for L2. In a separate 100-draw within-miRNA permutation of the trained scores, the corresponding 99th percentiles were 22 and 22.01. Both score-permutation comparisons favour the trained model; neither trained count strictly exceeds the random-weight 99th percentile. The locked rule required beating *both* nulls at $\alpha = 0.01$ for each claim. Thus both L1 and L2 fail. The result does not prove that the precise physical mechanism is site enumeration alone, but it does show that trained weights are unnecessary to reach the observed level of depletion under this pipeline. We withdraw the claim that this depletion is a new biological finding learned by DuplexCNN; the older observed comparisons remain descriptive and cannot rescue this gate. Full 100+100 arrays, input hashes, code and verdict are in `results/random_weight_falsify.json` and

results/recovered_inputs_20260927.json. The fresh TargetScan extraction and deterministic weight rebuild reproduce the archived training metric triple exactly; original ignored-input byte hashes were unavailable after a workspace loss, so byte identity with that earlier environment cannot be claimed. This negative result does not alter the separately specified CLIP primary endpoint; the refit-aware dependence interval is now computed and includes zero, and independent-library replication is still open.

6.2 Family-held-out transfer test and the pair-map decision

Before inspecting its outcome, we locked a transfer test over the five largest miRNA families in the deterministic 80,000-site TargetScan context++ build. For each family, we retrained the full DuplexCNN with context for six epochs on all other families, fit ridge+context ($\alpha = 10$) on the same training sites, and scored Pearson correlation on the held-out family. A fresh untrained DuplexCNN provided a random-weight comparison. Context normalization was fit on training families only. The locked rule required the CNN’s mean Pearson r across the five held-out families to exceed ridge+context *and* to exceed the random-weight mean by more than 0.05.

Held-out family	n	CNN r	Ridge+context r	Random-weight r
let-7-5p/98-5p	4,645	0.7162	0.7275	-0.1080
miR-15-5p group	4,143	0.7425	0.7949	-0.0231
miR-302-3p group	4,094	0.7127	0.7249	-0.1533
miR-17-5p group	4,029	0.7553	0.7396	-0.1517
miR-30-5p	2,748	0.5969	0.6434	-0.1195
Unweighted mean		0.7047	0.7261	-0.1111

The CNN beat ridge in only one of five families. Its unweighted mean exceeded the random-weight mean by 0.8159, passing that clause, but trailed ridge by 0.0214, failing the other. The joint locked verdict is *no support for transferable advantage over ridge*. This is TargetScan context++ transfer across families, not external CLIP validation. The test compares the *full* CNN with ridge, so it cannot isolate the pair-map branch. Its within-dataset gain was only +0.0029 Pearson r at $20\times$ training cost. We therefore deprecate the branch for new training without changing the archived artifact. Exact family outcomes: results/family_transfer.json; locked code: scripts/family_transfer.py.

6.3 Fixed-width absolute UTR-length-bin sensitivity audit

A coarse UTR-length-decile control in the earlier 30-miRNA held-out panel left a median-LOEUF contrast in 25 of 30 miRNAs. After the random-weight falsifier failed, we tested finer *absolute* UTR-length bins as a post-hoc sensitivity audit, not as a way to restore the trained-model claim. For each miRNA we selected the original 200 CNN-ranked candidates among LOEUF-mapped genes. A candidate was matchable only when its absolute length bin contained at least one non-candidate from the same miRNA universe; unmatchable candidates were excluded from both candidate and comparison medians. Three seeded control samples drew one gene per matchable candidate, with replacement inside its bin. Duplicate controls therefore reduce the effective comparison-set size, and there was no nearest-bin fallback.

Absolute bin width	Matchable/6,000	Median unique controls	Panels $\uparrow/\downarrow$	Descriptive p
25 nt	5,803	174	21/9	0.0214
50 nt	5,884	175	24/6	0.00072
100 nt	5,948	177	23/7	0.00261

An upward arrow means higher median LOEUF (less constraint) among the *matchable subset* of CNN-selected genes than the average matched control median. The *p* values are unadjusted one-sided sign tests, not pre-registered gates; the 25/50/100-nt audits reuse the same panel and cannot be counted as three independent replications. The loss of up to 197 candidates at 25 nt and sampling with replacement limit the interpretation. The result describes how a selected-gene pattern responds to finer length matching; because untrained networks matched the earlier depletion under the locked falsifier, it supplies no evidence that learned weights produced the pattern. The exact per-miRNA exclusions and medians are in the result JSON files for each width (`results/exact_length_bin_*bp.json`); the seeded drawing code is in `scripts/exact_length_bin_control.py`.

7 Falsification paths

miRTarBase 10.0 and TarBase v9.0 benchmarks are now done (see above); wet-lab: luciferase reporters for the top 10 diverse pairs; extended-panel CLIP rerun over the 132-miRNA ENCORI panel completed (delta +0.0060, *p*=0.048); second-resource TarBase AGO-CLIP label test passes the locked within-resource rule but is not independent-library proof; remaining: library-matched CLIP/RNA-seq normalization and independent-library replication.

References

- [Grimson et al.(2007)] Grimson A, et al. MicroRNA targeting specificity in mammals: determinants beyond seed pairing. *Mol Cell* 2007.
- [Agarwal et al.(2015)] Agarwal V, et al. Predicting effective microRNA target sites in mammalian mRNAs. *eLife* 2015.
- [Zhou et al.(2026)] Zhou K, et al. An encyclopedic regulatory and functional atlas of RNA interactomes (ENCORI). *Nat Methods* 2026.
- [Huang et al.(2022)] Huang HY, et al. miRTarBase update 2022: an information resource for experimentally validated miRNA–target interactions. *Nucleic Acids Res* 2022.
- [Skoufos et al.(2024)] Skoufos G, et al. TarBase-v9.0 extends experimentally supported miRNA-gene interactions to cell-types and virally encoded miRNAs. *Nucleic Acids Res* 2024.
- [Chen and Wang(2020)] Chen Y, Wang X. miRDB: an online database for prediction of functional microRNA targets. *Nucleic Acids Res* 2020.
- [Xin et al.(2016)] Xin J, et al. High-performance web services for querying gene and variant annotation. *Genome Biol* 2016.
- [Memczak et al.(2013)] Memczak S, et al. Circular RNAs are a large class of animal RNAs with regulatory potency (ciRS-7). *Nature* 2013.
- [Kolberg et al.(2023)] Kolberg L, et al. g:Profiler—interoperable web service for functional enrichment analysis and gene identifier mapping (2023 update). *Nucleic Acids Res* 2023.
- [Szklarczyk et al.(2023)] Szklarczyk D, et al. The STRING database in 2023: protein–protein association networks and functional enrichment analyses for any sequenced genome of interest. *Nucleic Acids Res* 2023.

- [ENCODE(2020)] ENCODE Project Consortium. Expanded encyclopaedias of DNA elements in the human and mouse genomes. *Nature* 2020.
- [Yates et al.(2015)] Yates A, et al. The Ensembl REST API: Ensembl data for any language. *Bioinformatics* 2015.
- [GTEx(2020)] GTEx Consortium. The GTEx Consortium atlas of genetic regulatory effects across human tissues. *Science* 2020.
- [Farh et al.(2005)] Farh KK-H, et al. The widespread impact of mammalian microRNAs on mRNA repression and evolution. *Science* 2005.
- [Kuleshov et al.(2016)] Kuleshov MV, et al. Enrichr: a comprehensive gene set enrichment analysis web server 2016 update. *Nucleic Acids Res* 2016.
- [Milacic et al.(2024)] Milacic M, et al. The Reactome Pathway Knowledgebase 2024. *Nucleic Acids Res* 2024.
- [RNAcentral(2021)] The RNAcentral Consortium. RNAcentral 2021: secondary structure integration, improved sequence search and new member databases. *Nucleic Acids Res* 2021.
- [Türei et al.(2021)] Türei D, et al. Integrated intra- and intercellular signaling knowledge for multicellular omics analysis. *Mol Syst Biol* 2021.
- [UniProt(2025)] The UniProt Consortium. UniProt: the Universal Protein Knowledgebase in 2025. *Nucleic Acids Res* 2025.
- [Karczewski et al.(2020)] K. J. Karczewski et al. The mutational constraint spectrum quantified from variation in 141,456 humans. *Nature* 581:434–443, 2020.
- [Rehm et al.(2015)] H. L. Rehm et al. ClinGen: the Clinical Genome Resource. *N Engl J Med* 372:2235–2242, 2015.
- [Collins et al.(2022)] R. L. Collins et al. A cross-disorder dosage sensitivity map of the human genome. *Cell* 185:3041–3055, 2022.
- [Köhler et al.(2021)] S. Köhler et al. The Human Phenotype Ontology in 2021. *Nucleic Acids Res* 49:D1207–D1217, 2021.
- [Landrum et al.(2018)] M. J. Landrum et al. ClinVar: improving access to variant interpretations and supporting evidence. *Nucleic Acids Res* 46:D1062–D1067, 2018.
- [Tsherniak et al.(2017)] A. Tsherniak et al. Defining a cancer dependency map. *Cell* 170:564–576, 2017.
- [Meldal et al.(2019)] B. H. M. Meldal et al. Complex Portal 2018: extended content and enhanced visualization tools for macromolecular complexes. *Nucleic Acids Res* 47:D550–D558, 2019.
- [Brown et al.(2015)] G. R. Brown et al. Gene: a gene-centered information resource at NCBI. *Nucleic Acids Res* 43:D36–D42, 2015.
- [Huang et al.(2010)] N. Huang et al. Characterising and predicting haploinsufficiency in the human genome. *PLoS Genet* 6:e1001154, 2010.
- [Zeng et al.(2024)] T. Zeng et al. Bayesian estimation of gene constraint from an evolutionary model with gene features. *Nat Genet* 56:1632–1643, 2024.

[Kinsella et al.(2011)] R. J. Kinsella et al. Ensembl BioMarts: a hub for data retrieval across taxonomic space. *Database* 2011:bar030, 2011.

[Seal et al.(2023)] R. L. Seal et al. Genenames.org: the HGNC resources in 2023. *Nucleic Acids Res.* 51:D1003–D1009, 2023.

[Orphanet(2024)] Orphanet. Orphadata: free access data from Orphanet, product 6 (genes associated with rare diseases). INSERM, www.orphadata.com, accessed 2026.

[Baldarelli et al.(2024)] R. M. Baldarelli et al. Mouse Genome Informatics: an integrated knowledgebase system for the laboratory mouse. *Genetics* 227:iyae031, 2024.

8 Mathematical derivations

Notation matches `src/duplexcnn` and `scripts/clip_falsify.py`; every statistic below appears in a committed results JSON.

8.1 Convolutional duplex scoring

The pair-map P is scored by stacked convolutions. One layer with kernel $W \in \mathbb{R}^{C_{out} \times C_{in} \times k}$ and bias b :

$$h_{c,j}^{(\ell+1)} = \sigma \left(\sum_{c'} \sum_{u=0}^{k-1} W_{c,c',u} h_{c',j+u-p}^{(\ell)} + b_c \right), \quad \sigma(x) = \max(0, x), \quad (3)$$

with padding $p = \lfloor k/2 \rfloor$ keeping length fixed so positional pairing information survives to the head. Training minimizes MSE against the TargetScan context++ score; the held-out unit is the *gene*, so no transcript can contribute to both fit and evaluation.

8.2 Pearson and Spearman benchmarks

Held-out agreement with context++ is

$$r = \frac{\sum_i (y_i - \bar{y})(\hat{y}_i - \bar{\hat{y}})}{\sqrt{\sum_i (y_i - \bar{y})^2 \sum_i (\hat{y}_i - \bar{\hat{y}})^2}}, \quad r_s = r(\text{rank}(y), \text{rank}(\hat{y})), \quad (4)$$

the second being Pearson on ranks, robust to monotone distortion of the score scale. DuplexCNN: $r = 0.8121$, $r_s = 0.7928$ (committed ablation JSON).

8.3 Fisher exact test for CLIP enrichment

For one miRNA, genes partition by CNN min-score (high/low) and CLIP support (supported/not) into a 2×2 table with margins fixed by construction. Under the null of independence the count x of high-scoring CLIP-supported genes is hypergeometric:

$$P(X = x) = \frac{\binom{K}{x} \binom{N-K}{n-x}}{\binom{N}{n}}, \quad p = \sum_{x' \geq x} P(X = x'), \quad (5)$$

an exact small-sample test that needs no asymptotic approximation.

8.4 Odds ratio

Effect size is reported as

$$\text{OR} = \frac{x/(n-x)}{(K-x)/(N-K-n+x)}, \quad (6)$$

so a per-miRNA AUROC of 0.55 can be read directly as an enrichment strength rather than a bare significance flag.

8.5 AUROC as a Mann–Whitney statistic

The per-miRNA CNN-min AUROC against CLIP labels is

$$\text{AUROC} = \frac{1}{n_+n_-} \sum_{i \in +} \sum_{j \in -} \left(\mathbf{1}[s_i > s_j] + \frac{1}{2} \mathbf{1}[s_i = s_j] \right), \quad (7)$$

i.e. the probability a random CLIP-supported gene outscores a random unsupported one; chance is 0.5 by construction.

8.6 Bootstrap confidence interval for ΔAUROC

The base-vs-base+CNN comparison resamples (miRNA, gene) units with replacement $B = 1000$ times, refits nothing (scores are fixed), and recomputes the AUROC gap:

$$\Delta_b = \text{AUROC}_b^{\text{base+CNN}} - \text{AUROC}_b^{\text{base}}, \quad \text{CI}_{95} = [\hat{Q}_\Delta(0.025), \hat{Q}_\Delta(0.975)], \quad (8)$$

which yielded $[+0.0057, +0.0093]$ on the 23-miRNA panel — the interval excludes zero, so the sign of the effect is not sampling noise.

8.7 Permutation null

To test whether the CNN adds signal beyond covariates within miRNA, labels are shuffled *inside* each miRNA stratum $B = 20$ times:

$$p = \frac{1 + \#\{b : \Delta_b \geq \Delta_{\text{obs}}\}}{B + 1}, \quad (9)$$

preserving miRNA-level base rates under the null. The 23-miRNA panel gave $p = 0.048$; the 136-miRNA rerun (this paper’s extended falsification) uses the identical null so the two numbers are commensurable.

8.8 Seed-match enumerator

The discovery scan counts exact seed complementarity. For UTR u and miRNA m with seed σ_m (positions 2–8),

$$n_{8\text{mer}}(u, m) = \sum_j \mathbf{1}[u_{j:j+8} = \text{revcomp}(\sigma_m A1)], \quad n_{\text{total}} = \sum_m \sum_u (n_{8\text{mer}} + n_{7\text{mer-A1}}), \quad (10)$$

enumerating 1,453,439 novel non-conserved matches — the scale that makes a learned prioritizer necessary and the falsification non-trivial.

8.9 Wilcoxon signed-rank for paired per-gene comparisons

Where designs are compared per gene g (full vs ablated model), differences d_g are ranked by magnitude and

$$W = \sum_{g:d_g>0} \text{rank}(|d_g|), \quad p \text{ from the exact null distribution of } W, \quad (11)$$

a distribution-free paired test whose $W = 0$ (all differences one direction) is the strongest possible ordinal outcome at a given n .